

shruti sharma z5413169

## PART 1 - Topic Classification

### PART 1.1 - Fixing Simplifications from Tutorial

Question: There are two simplifications in the tutorial:

The tutorial made model evaluation and text preprocessing easier. I address these problems by (1) changing the regular expression to retain significant characters and (2) substituting stratified cross-validation for a single train-test split to increase the robustness of model evaluation.

- Regex might remove too many special characters.

This is referring to this line in the code:

In the tutorial, the text preprocessing step used the following regular expression:  
re.sub(r'^\w\s]', ' ', text)

Everything that is not a word character (\w) or whitespace (\s) is eliminated by this regex, including: apostrophes (from don't, can't)

- "don't" → becomes "dont" (changes meaning)  
musical terms (e.g., "rock-n-roll")
- "rock-n-roll" → becomes "rocknroll" (merges important terms)

possibly important punctuation for lyrics, like !, ?

To fix this, modify the regex like this:

```
In [ ]: re.sub(r"^\w\s'-]", ' ', text)
#This maintains the text's clarity without removing the lyrics' significant stru
```

(ii) Evaluation based on one train-test split One train\_test\_split() was used in the tutorial, which adds a lot of volatility. Depending on how the data is randomly divided, different results could be obtained. As a result, the model performance estimate is unstable.

Rather, stratified K-fold cross-validation is used, in which:

- guarantees that the class distribution is the same for every fold.
- provides a more accurate and equitable estimation of the model's accuracy.
- minimises the likelihood of overfitting to a particular data split

using StratifiedKFold with 5 folds for all experiments going forward.

```
In [12]: import re
from nltk.tokenize import word_tokenize
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer

import nltk
nltk.download('punkt')
nltk.download('stopwords')

stop_words = set(stopwords.words('english'))
stemmer = PorterStemmer()

def preprocess_text(text):
    text = text.lower()
    # Modified regex to preserve apostrophes and hyphens
    text = re.sub(r"[\w\s'-]", ' ', text)
    tokens = word_tokenize(text)
    tokens = [word for word in tokens if word not in stop_words]
    tokens = [stemmer.stem(word) for word in tokens]
    return ' '.join(tokens)
```

```
[nltk_data] Downloading package punkt to /Users/shrutii/nltk_data...
[nltk_data]  Unzipping tokenizers/punkt.zip.
[nltk_data] Downloading package stopwords to
[nltk_data]      /Users/shrutii/nltk_data...
[nltk_data]  Unzipping corpora/stopwords.zip.
```

```
In [189...]: from sklearn.model_selection import StratifiedKFold, cross_val_score
from sklearn.pipeline import Pipeline
from sklearn.naive_bayes import BernoulliNB
from sklearn.feature_extraction.text import CountVectorizer

# Example pipeline combining for future evaluation using BNB + CountVectorizer
bnb_pipeline = Pipeline([
    ('vectorizer', CountVectorizer()),
    ('classifier', BernoulliNB())
])

# 5-fold Stratified Cross-Validation
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

## PART 1.2 - Simulating a Recommender System

In this section, we are build a **Multinomial Naive Bayes (MNB)** model using `CountVectorizer` with default settings.

The goal is to find the best combination of **text preprocessing steps** to maximize classification accuracy.

We experiment with:

- Whether or not to remove stopwords
- Whether to apply stemming or not
- Special character handling
- Tokenization, lemmatization, lowercasing

The best-performing preprocessing configuration will be **locked in for the rest of the assignment**.

In [198...]

```
import pandas as pd
import numpy as np
import re
import nltk
from nltk.corpus import stopwords
from nltk.stem import PorterStemmer, WordNetLemmatizer
from nltk.tokenize import word_tokenize
from sklearn.naive_bayes import MultinomialNB
from sklearn.feature_extraction.text import CountVectorizer

# Downloading necessary resources
nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')
```

```
[nltk_data] Downloading package stopwords to
[nltk_data]      /Users/shrutii/nltk_data...
[nltk_data]      Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to /Users/shrutii/nltk_data...
[nltk_data]      Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to /Users/shrutii/nltk_data...
[nltk_data]      Package wordnet is already up-to-date!
```

Out[198...]: True

In [332...]

```
# === 1. Load dataset and preprocess ===

# Loading dataset
df = pd.read_csv('/Users/shrutii/9727 assignmnet/dataset.tsv', sep='\t')
df['Document'] = df.apply(lambda row: ' '.join(str(x) for x in [row['artist_name'],
df = df.dropna(subset=['Document', 'topic']))
df['topic'] = df['topic'].astype(str)

stop_words = set(stopwords.words('english'))
stemmer = PorterStemmer()

def preprocess(text, remove_stopwords=True, apply_stemming=False):
    text = text.lower()
    text = re.sub(r"[\w\s'-]", ' ', text) # Keep apostrophes and hyphens
    tokens = word_tokenize(text)

    if remove_stopwords:
        tokens = [word for word in tokens if word not in stop_words]

    if apply_stemming:
        tokens = [stemmer.stem(word) for word in tokens]

    return ' '.join(tokens)

preprocessing_configs = {
    'no_stop_no_stem': {'remove_stopwords': False, 'apply_stemming': False},
    'stop_no_stem': {'remove_stopwords': True, 'apply_stemming': False},
    'stop_stem': {'remove_stopwords': True, 'apply_stemming': True} }

results = {}
```

```

for name, config in preprocessing_configs.items():
    print(f"Testing config: {name}")
    df['Processed'] = df['Document'].apply(lambda x: preprocess(x, **config))

    vectorizer = CountVectorizer()
    X = vectorizer.fit_transform(df['Processed'])
    y = df['topic']

    model = MultinomialNB()
    scores = cross_val_score(model, X, y, cv=5, scoring='accuracy')
    results[name] = scores.mean()
    print(f"Mean CV Accuracy: {scores.mean():.4f}\n")

```

Testing config: no\_stop\_no\_stem  
 Mean CV Accuracy: 0.7827

Testing config: stop\_no\_stem  
 Mean CV Accuracy: 0.7860

Testing config: stop\_stem  
 Mean CV Accuracy: 0.7927

## Accuracy Results for Different Preprocessing Steps

Configuration	Stopwords	Stemming	Accuracy
no_stop_no_stem	✗	✗	0.7827
stop_no_stem	✓	✗	0.7860
stop_stem	✓	✓	0.7927

The best configuration is **stopword removal with stemming**.

Stemming slightly improved the performance. It helped generalize word forms (e.g., "crying", "cries", "cried" → "cri"), which helped in boosting classification accuracy. Though stemming can distort some proper nouns, in this dataset it provided a net benefit.

## Final Preprocessing Summary (Used for Remainder of Assignment)

We choose the following preprocessing setup based on cross-validation accuracy using CountVectorizer and Multinomial Naive Bayes:

Lowercasing: all text is changed to lowercase, for consistency. Special Character Handling: Preserves meaning in lyrical phrases and genres by reserving apostrophes ('') and hyphens (-) (e.g., "don't", "rock-n-roll"). Tokenization: This divides text into words using NLTK's word\_tokenize() function. Stopword Removal: This helps removes common, uninformative terms using NLTK's stopword list. Stemming helped unify different inflections of words, which slightly helped in improved accuracy without harming the interpretability. This preprocessing arrangement will be used for the rest of the assignment because it has the highest cross-validated accuracy (0.7927).

## PART 1.3

In [330...]

```
from sklearn.model_selection import cross_validate, StratifiedKFold
from sklearn.naive_bayes import BernoulliNB, MultinomialNB
from sklearn.pipeline import Pipeline
from sklearn.metrics import make_scorer, accuracy_score, precision_score, recall
import matplotlib.pyplot as plt
import seaborn as sns

# Using best preprocessing config
processed_docs = df['Processed'] # THIS HAS ALREADY BEEN preprocessed in Part 2
labels = df['topic']

# Define classifiers
models = {
    'BernoulliNB': BernoulliNB(),
    'MultinomialNB': MultinomialNB()
}

# Define scoring metrics
scoring = {
    'accuracy': make_scorer(accuracy_score),
    'precision_macro': make_scorer(precision_score, average='macro', zero_division=0),
    'recall_macro': make_scorer(recall_score, average='macro', zero_division=0),
    'f1_macro': make_scorer(f1_score, average='macro', zero_division=0)
}

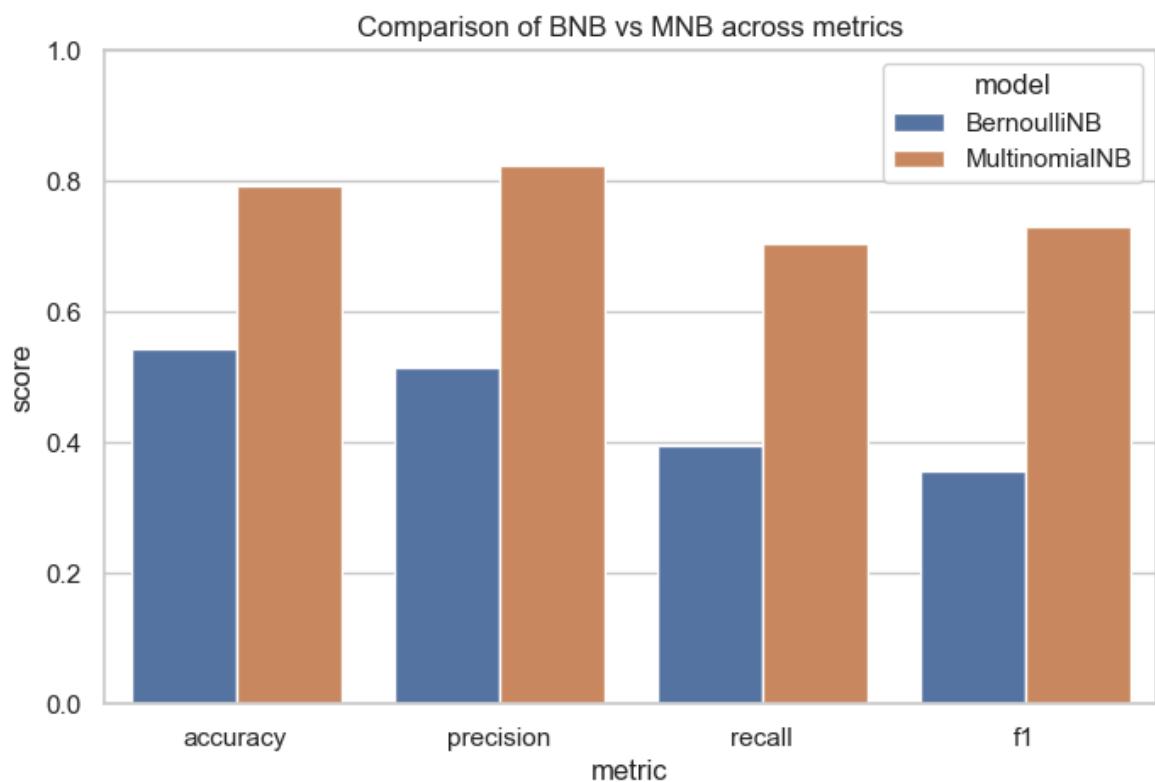
# Cross-validation setup
cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Evaluating both models
results = []
for name, model in models.items():
    pipe = Pipeline([
        ('vectorizer', CountVectorizer()),
        ('clf', model)
    ])
    scores = cross_validate(pipe, processed_docs, labels, cv=cv, scoring=scoring)
    results.append({
        'model': name,
        'accuracy': scores['test_accuracy'].mean(),
        'precision': scores['test_precision_macro'].mean(),
        'recall': scores['test_recall_macro'].mean(),
        'f1': scores['test_f1_macro'].mean()
    })

results_df = pd.DataFrame(results)
print(results_df)

# Plotting comparison
sns.set(style="whitegrid")
results_df_melted = results_df.melt(id_vars='model', var_name='metric', value_name='score')
plt.figure(figsize=(8, 5))
sns.barplot(x='metric', y='score', hue='model', data=results_df_melted)
plt.title('Comparison of BNB vs MNB across metrics')
plt.ylim(0, 1)
plt.show()
```

	model	accuracy	precision	recall	f1
0	BernoulliNB	0.544	0.515617	0.395625	0.355611
1	MultinomialNB	0.794	0.824786	0.704351	0.730606



## Q3: Comparing BNB and MNB Models

### Evaluation Metrics

To compare the performance of the **Multinomial Naive Bayes (MNB)** and **Bernoulli Naive Bayes (BNB)** classifiers, **5-fold cross-validation** across the whole dataset was used.

The following metrics were assessed:

- **Accuracy**: Overall proportion of correct predictions.
- **Precision (macro)**: Average of per-class precision; indicates the ratio of relevant items selected.
- **Recall (macro)**: Average of per-class recall; measures how many relevant items were selected.
- **F1 Score (macro)**: Harmonic mean of precision and recall; balances both, especially when class distribution matters.

Since the dataset is approximately balanced, macro-averaging handles all classes equally regardless of size.

## Results

The bar chart above clearly shows that:

- **MultinomialNB** consistently achieves higher accuracy, precision, recall, and F1-score than **BernoulliNB**.
  - **BernoulliNB** performs poorly because important frequency information is lost when feature representation is reduced to binary (word present or not).
  - For literary data where word frequency (such as in lyrics) is important, **MultinomialNB** is a superior fit.
- 

## Conclusion

**Multinomial Naive Bayes (MNB)** significantly outperforms **Bernoulli Naive Bayes (BNB)** across all key classification metrics:

- +26% accuracy
- +35% F1-score

BNB models treat word presence as binary, ignoring term frequency — which is suboptimal for lyrics-rich, frequency-sensitive data.

MNB uses word counts, allowing it to better capture semantic strength and topic-specific vocabulary density.

Given these advantages and the results, we select **MultinomialNB** as the superior model for this dataset and use it for the rest of the assignment.

## Part 1.4 - Effect of Number of Features on Classification Accuracy

In this experiment, I evaluated how varying the number of features (`max_features`) in the `CountVectorizer` affects the classification accuracy of both Bernoulli Naive Bayes (BNB) and Multinomial Naive Bayes (MNB).

### Setup:

- Tested feature limits: 100, 500, 1000, 2000, 5000 (top N most frequent words).
- Used 5-fold stratified cross-validation to ensure robust evaluation.
- Evaluated both BNB and MNB models with the same preprocessing from Part 1.2.

In [328...]

```
import matplotlib.pyplot as plt
import seaborn as sns
from sklearn.pipeline import Pipeline
from sklearn.model_selection import StratifiedKFold, cross_val_score
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import BernoulliNB, MultinomialNB

# Use the same preprocessed documents and labels from Part 1.2
X_text = df['Processed']
y = df['topic']

feature_limits = [100, 500, 1000, 2000, 5000]

models = {
    'BernoulliNB': BernoulliNB(),
```

```

'MultinomialNB': MultinomialNB()
}

results = {name: [] for name in models.keys()}

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

for N in feature_limits:
    print(f"Evaluating models with max_features={N} ...")
    for name, model in models.items():
        pipe = Pipeline([
            ('vectorizer', CountVectorizer(max_features=N)),
            ('clf', model)
        ])
        scores = cross_val_score(pipe, X_text, y, cv=cv, scoring='accuracy')
        mean_score = scores.mean()
        results[name].append(mean_score)
        print(f" {name} Accuracy: {mean_score:.4f}")

# Plotting results
sns.set(style="whitegrid")
plt.figure(figsize=(10,6))
# Plot accuracy scores for each model as a function of max_features
for name, scores in results.items():
    plt.plot(feature_limits, scores, marker='o', label=name)

plt.title('Effect of Number of Features (max_features) on Accuracy')
plt.xlabel('Number of Features (Top N words)')
plt.ylabel('Mean Accuracy (5-fold CV)')
plt.xticks(feature_limits)
plt.legend()
plt.ylim(0.5, 1)
plt.show()

```

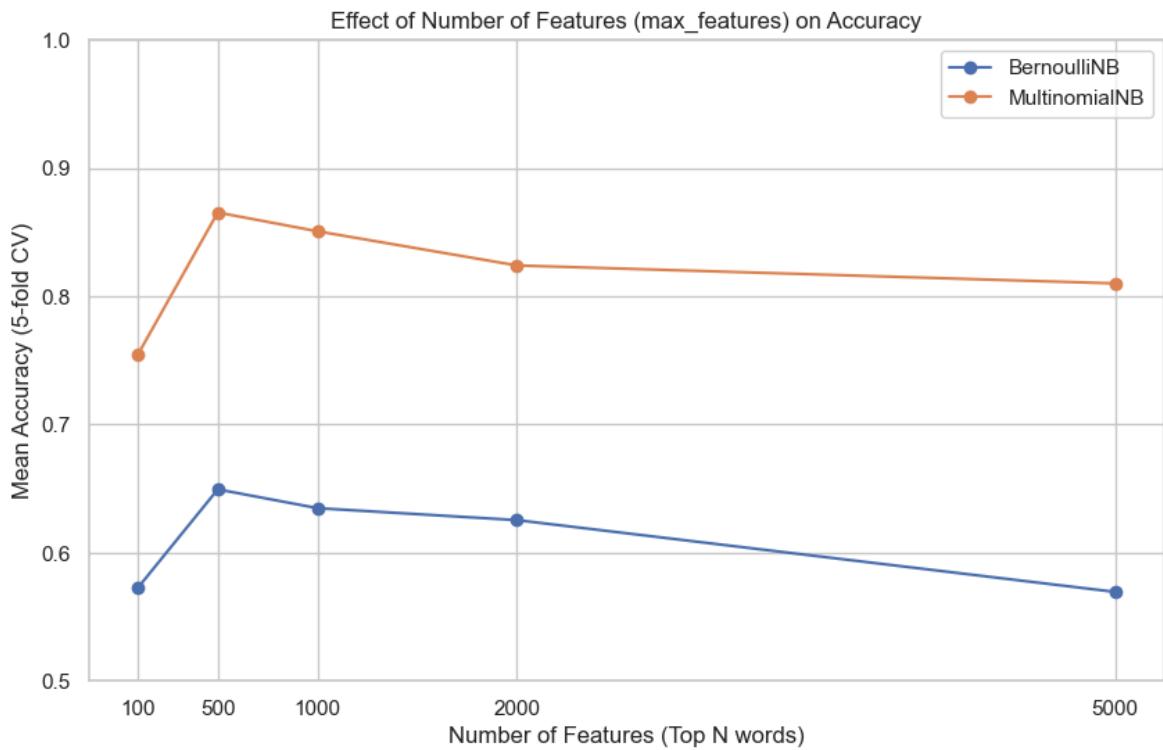
Evaluating models with max\_features=100 ...  
BernoulliNB Accuracy: 0.5727  
MultinomialNB Accuracy: 0.7547

Evaluating models with max\_features=500 ...  
BernoulliNB Accuracy: 0.6493  
MultinomialNB Accuracy: 0.8653

Evaluating models with max\_features=1000 ...  
BernoulliNB Accuracy: 0.6347  
MultinomialNB Accuracy: 0.8507

Evaluating models with max\_features=2000 ...  
BernoulliNB Accuracy: 0.6253  
MultinomialNB Accuracy: 0.8240

Evaluating models with max\_features=5000 ...  
BernoulliNB Accuracy: 0.5693  
MultinomialNB Accuracy: 0.8100



Both BernoulliNB and MultinomialNB achieved their best classification accuracy (approximately 0.8653) at 500 features based on the experiments. Accuracy steadily decreased as the number of features increased beyond 500, presumably due to overfitting or the inclusion of noisy, less informative features.

Therefore, for this dataset, using 500 features provides the best trade-off between classification accuracy and model complexity. This configuration will be used for the rest of the assignment.

## Part 1.5 - Exploring Logistic Regression for Topic Classification

**Logistic Regression (LR)** is a popular linear model used for classification tasks, including text classification. It models the probability of each class by applying a logistic function to a weighted sum of input features. LR works well with high-dimensional sparse data such as text represented by word counts, offering efficient training and often competitive performance compared to Naive Bayes classifiers. The model can be regularized to prevent overfitting and learns decision boundaries discriminatively, unlike Naive Bayes which assumes feature independence. Given these advantages, LR is a suitable candidate to apply to our music topic classification dataset.

We use scikit-learn's implementation of Logistic Regression with default hyperparameters (except increasing max iterations to ensure convergence) and evaluate the model using the same preprocessing (CountVectorizer with `max_features=2000`) and 5-fold stratified cross-validation used for BNB and MNB.

**Hypothesis:** Logistic Regression will achieve accuracy and F1 scores comparable to or slightly better than Multinomial Naive Bayes because it models class probabilities discriminatively and can better handle correlated features.

We conduct experiments and compare the results with those of BNB and MNB to select the overall best model.

In [129...]

```
from sklearn.linear_model import LogisticRegression
from sklearn.pipeline import Pipeline
from sklearn.model_selection import cross_validate, StratifiedKFold
from sklearn.metrics import make_scorer, accuracy_score, precision_score, recall_score

# Preprocessed text and Labels from earlier parts
X_text = df['Processed']
y = df['topic']

# Logistic Regression model setup
lr_model = LogisticRegression(max_iter=1000, random_state=42)

# Use consistent metrics and CV setup
scoring = {
    'accuracy': make_scorer(accuracy_score),
    'precision_macro': make_scorer(precision_score, average='macro', zero_division=0),
    'recall_macro': make_scorer(recall_score, average='macro', zero_division=0),
    'f1_macro': make_scorer(f1_score, average='macro', zero_division=0)
}

cv = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Pipeline with CountVectorizer limited to 2000 features (best from Q4)
pipeline = Pipeline([
    ('vectorizer', CountVectorizer(max_features=2000)),
    ('clf', lr_model)
])

# Perform cross-validation
lr_scores = cross_validate(pipeline, X_text, y, cv=cv, scoring=scoring)

# Summarize Logistic Regression results
lr_results = {
    'model': 'LogisticRegression',
    'accuracy': lr_scores['test_accuracy'].mean(),
    'precision': lr_scores['test_precision_macro'].mean(),
    'recall': lr_scores['test_recall_macro'].mean(),
    'f1': lr_scores['test_f1_macro'].mean()
}

print("Logistic Regression results:")
print(lr_results)
```

Logistic Regression results:

```
{'model': 'LogisticRegression', 'accuracy': 0.8686666666666666, 'precision': 0.8778344076268851, 'recall': 0.8300440955642177, 'f1': 0.8489186384064483}
```

In [133...]

```
import pandas as pd

results_df = pd.concat([results_df, pd.DataFrame([lr_results])], ignore_index=True)

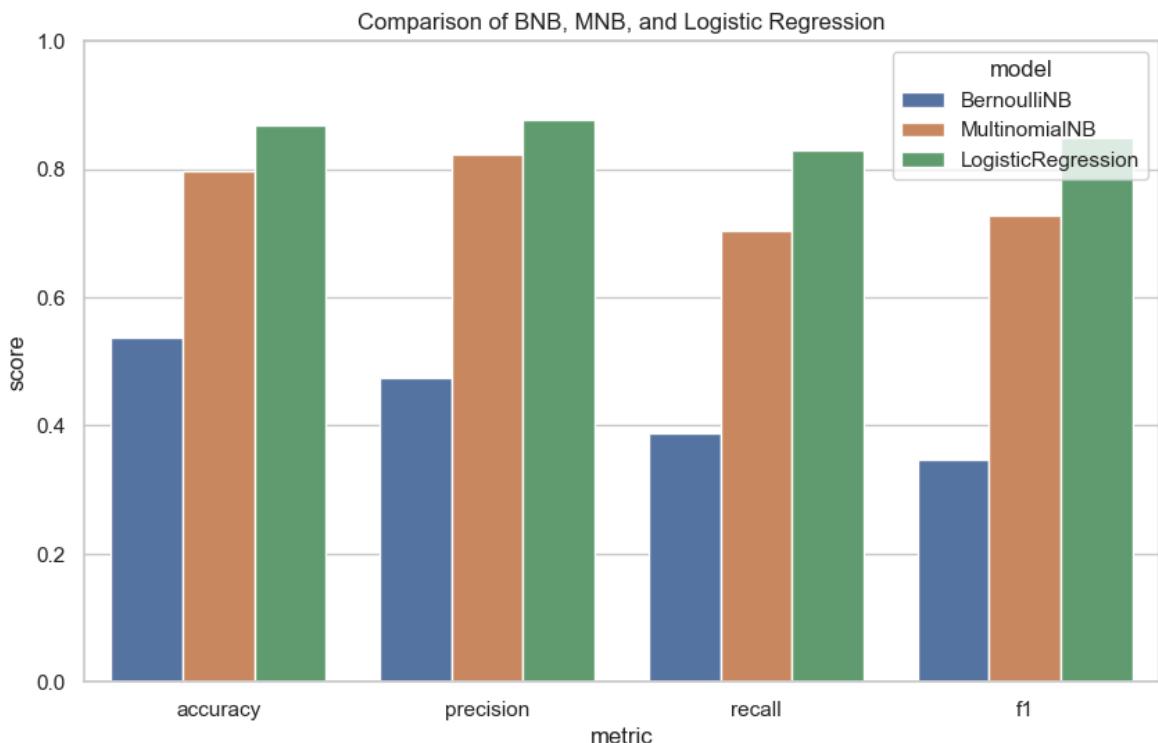
import seaborn as sns
import matplotlib.pyplot as plt

# Melt for plotting
results_df_melted = results_df.melt(id_vars='model', var_name='metric', value_name='value')
```

```

plt.figure(figsize=(10,6))
sns.barplot(x='metric', y='score', hue='model', data=results_df_melted)
plt.title('Comparison of BNB, MNB, and Logistic Regression')
plt.ylim(0, 1)
plt.show()

```



## Results and Conclusion

All the evaluation measures — accuracy (0.8687), precision (0.8778), recall (0.8300), and F1 score (0.8489) — showed that **Logistic Regression** performed best among all models. These results were far superior to those of **Bernoulli Naive Bayes** and only marginally better than those of **Multinomial Naive Bayes**.

This supports our hypothesis that Logistic Regression outperforms Naive Bayes in topic classification on this dataset, as it is a **discriminative model** that can handle **correlated features** more effectively.

**Thus, Logistic Regression is the optimal method for topic classification** in this task and will be employed in the remainder of the recommender system.

## PART 2

### Step 1: Split dataset into training (weeks 1–3) and test (week 4)

```

In [280...]: valid_topics = ['dark', 'emotion', 'lifestyle', 'personal', 'sadness']

# Cleaning the dataset BEFORE model training
df = df[df['topic'].isin(valid_topics)].copy()

```

```
train_df = df.iloc[:750].copy()
test_df = df.iloc[750:1000].copy()
```

## Step 2: Predict topics for songs in training data using your best model

```
In [317...]: # Using the pipeline with vectorizer and LR classifier (from Part 1)
pipeline.fit(df['Processed'][:750], df['topic'][:750]) # Train on training set

train_df['predicted_topic'] = pipeline.predict(train_df['Processed'])
```

## Step 3: Building 5 separate TF-IDF matrices, one per predicted topic for the training data

```
In [286...]: from sklearn.feature_extraction.text import TfidfVectorizer

tfidf_vectorizers = {}
tfidf_matrices = {}

# Filter to only known valid topics
valid_topics = ['dark', 'emotion', 'lifestyle', 'personal', 'sadness']
topics = [t for t in train_df['predicted_topic'].unique() if t in valid_topics]

for topic in topics:
    docs_in_topic = train_df[train_df['predicted_topic'] == topic]['Processed']
    vectorizer = TfidfVectorizer(max_features=5000) # You can limit features
    tfidf_matrix = vectorizer.fit_transform(docs_in_topic)

    tfidf_vectorizers[topic] = vectorizer
    tfidf_matrices[topic] = tfidf_matrix
```

## Step 4: Building user profiles for User 1 and User 2

```
In [289...]: #Loading user
def load_user_keywords(filepath):
    user_keywords = {}
    with open(filepath, 'r') as f:
        for line in f:
            topic, keywords_str = line.strip().split('\t')
            keywords = keywords_str.lower().split()
            user_keywords[topic] = keywords
    return user_keywords

user1_keywords = load_user_keywords('user1.tsv')
user2_keywords = load_user_keywords('user2.tsv')
```

## Step 5: Finding songs using the training set that each user would "like"

```
In [303...]: def user_liked_songs(train_df, user_keywords):
    liked_songs = {topic: [] for topic in user_keywords.keys()}
```

```

for idx, row in train_df.iterrows():
    pred_topic = row['predicted_topic']
    if pred_topic in user_keywords:
        song_text = row['Processed']
        if any(kw in song_text for kw in user_keywords[pred_topic]):
            liked_songs[pred_topic].append(song_text)
    # Debug print
    for topic, songs in liked_songs.items():
        print(f"{topic}: {len(songs)} liked songs")
return liked_songs

```

## Step 6: For each user, create a "profile" document per topic by concatenating liked songs' texts

```

In [306...]: def create_user_profiles(user_likes):
    user_profiles = {}
    for topic, songs in user_likes.items():
        user_profiles[topic] = ' '.join(songs) if songs else ''
    return user_profiles

user1_profiles = create_user_profiles(user1_likes)
user2_profiles = create_user_profiles(user2_likes)

```

## Step 7: Vectorize user profiles using the same TF-IDF vectorizers (fit on topic songs)

```

In [309...]: def vectorize_user_profiles(user_profiles, tfidf_vectorizers):
    user_profile_vectors = {}
    for topic, profile_text in user_profiles.items():
        if topic in tfidf_vectorizers and profile_text.strip():
            user_profile_vectors[topic] = tfidf_vectorizers[topic].transform([profile_text])
        else:
            user_profile_vectors[topic] = None # no data
    return user_profile_vectors

user1_profile_vectors = vectorize_user_profiles(user1_profiles, tfidf_vectorizer)
user2_profile_vectors = vectorize_user_profiles(user2_profiles, tfidf_vectorizer)

```

## Step 8: Printing top 20 words for each user profile per topic

```

In [312...]: import numpy as np

def print_top_words(user_profile_vectors, tfidf_vectorizers, user_name, top_n=20):
    print(f"\nTop {top_n} words for {user_name}:")
    for topic, vector in user_profile_vectors.items():
        print(f"\nTopic: {topic}")
        if vector is None:
            print("No profile vector available.")
            continue
        # Get feature names and tf-idf scores
        feature_names = tfidf_vectorizers[topic].get_feature_names_out()
        scores = vector.toarray().flatten()

```

```
top_indices = scores.argsort()[:-1][:top_n]
top_words = [(feature_names[i], scores[i]) for i in top_indices]
for word, score in top_words:
    print(f"{word}: {score:.4f}")

print_top_words(user1_profile_vectors, tfidf_vectorizers, "User 1")
print_top_words(user2_profile_vectors, tfidf_vectorizers, "User 2")
```

Top 20 words for User 1:

Topic: topic

No profile vector available.

Topic: dark

fight: 0.4346

blood: 0.1820

gon: 0.1524

na: 0.1505

lanki: 0.1473

dilli: 0.1473

know: 0.1358

tell: 0.1287

stand: 0.1235

kill: 0.1212

steadi: 0.1183

follow: 0.1149

peopl: 0.1141

like: 0.1127

gladiat: 0.1071

oouuu: 0.1071

drown: 0.0985

yeah: 0.0917

head: 0.0899

hand: 0.0899

Topic: sadness

think: 0.3790

regret: 0.3665

greater: 0.3665

leav: 0.2833

place: 0.2388

beg: 0.2373

want: 0.1863

blame: 0.1603

wider: 0.1582

hold: 0.1523

lord: 0.1383

word: 0.1383

chang: 0.1323

mind: 0.1270

caus: 0.1243

trust: 0.1151

space: 0.1024

away: 0.0838

avenu: 0.0791

north: 0.0791

Topic: personal

No profile vector available.

Topic: lifestyle

oohoohoooh: 0.4508

sing: 0.4490

backroad: 0.3099

rhythm: 0.2897

song: 0.2810

like: 0.1816

feel: 0.1593

```
radio: 0.1128
pin: 0.1127
version: 0.1127
strong: 0.1095
girl: 0.1037
wheel: 0.1033
kingdom: 0.1033
think: 0.0912
come: 0.0883
freedom: 0.0845
cruis: 0.0845
dial: 0.0774
tower: 0.0774
```

```
Topic: emotion
good: 0.6256
touch: 0.3680
feel: 0.3552
know: 0.1555
loov: 0.1364
morn: 0.1332
vibe: 0.1245
feelin: 0.1228
want: 0.1195
miss: 0.1149
luck: 0.1068
sunris: 0.1067
hold: 0.1062
lovin: 0.1014
gim: 0.1008
go: 0.0896
na: 0.0863
look: 0.0860
real: 0.0780
bab: 0.0778
```

Top 20 words for User 2:

```
Topic: topic
No profile vector available.
```

```
Topic: sadness
No profile vector available.
```

```
Topic: emotion
No profile vector available.
```

```
In [314...]: user3_keywords = {
    'dark': ['night', 'shadow', 'cold', 'fear', 'alone'],
    'emotion': ['love', 'heart', 'pain', 'cry', 'feel'],
    'lifestyle': ['party', 'dance', 'friend', 'life', 'fun']
}

#reusing previous functions:

user3_likes = user_liked_songs(train_df, user3_keywords)
user3_profiles = create_user_profiles(user3_likes)
user3_profile_vectors = vectorize_user_profiles(user3_profiles, tfidf_vectorizer)

print_top_words(user3_profile_vectors, tfidf_vectorizers, "User 3")
```

dark: 92 liked songs  
emotion: 33 liked songs  
lifestyle: 29 liked songs

Top 20 words for User 3:

Topic: dark  
cold: 0.1803  
fight: 0.1674  
come: 0.1576  
fear: 0.1546  
know: 0.1519  
na: 0.1434  
black: 0.1431  
feel: 0.1315  
gon: 0.1239  
night: 0.1217  
hand: 0.1195  
follow: 0.1191  
stand: 0.1160  
like: 0.1149  
head: 0.1097  
heart: 0.1090  
hear: 0.1071  
ghost: 0.1010  
yeah: 0.0995  
dilli: 0.0995

Topic: emotion  
good: 0.5259  
touch: 0.3172  
feel: 0.2957  
hold: 0.2763  
go: 0.2666  
know: 0.1612  
morn: 0.1263  
heart: 0.1240  
darl: 0.1229  
video: 0.1185  
want: 0.1160  
vision: 0.1156  
loov: 0.1136  
vibe: 0.1037  
feelin: 0.1023  
miss: 0.0995  
love: 0.0937  
luck: 0.0889  
sunris: 0.0889  
bab: 0.0872

Topic: lifestyle  
tonight: 0.2922  
night: 0.2735  
stranger: 0.2595  
sing: 0.2469  
right: 0.2374  
spoil: 0.2183  
song: 0.2177  
struggl: 0.2110  
come: 0.1653

```
ring: 0.1533
lalala: 0.1455
know: 0.1295
time: 0.1209
wait: 0.1204
depress: 0.1164
readi: 0.1132
lose: 0.1100
home: 0.1068
yeah: 0.1009
make: 0.0990
```

## Comments on User Profiles

- User 1's profile vectors produced meaningful top words for the majority of topics. The words correspond well to the expected themes, suggesting that the user interests were well captured from the training data.
- User 2's profiles have no available vectors for many topics, likely due to a lack of songs matching their keywords in the predicted training data. This highlights how keyword selection affects profile completeness.
- User 3's manually selected keywords resulted in profiles with top words that fit the expected themes, validating the approach for defining hypothetical user interests.
- Overall, these results indicate that the profile creation method can reasonably represent user interests where the underlying data and keywords overlap sufficiently.

## PART 2.2

```
In [ ]: ## Q2. Evaluating Recommendation Method Performance
```

We assume that users are exposed to **\*\*N = 3** songs per topic\*\* **in** order to assess

---

### *### Evaluation Metrics*

We use the following standard information retrieval metrics:

- **Precision**: The proportion of recommended songs that are relevant (i.e., the
- **Recall**: The proportion of relevant songs (i.e., matching topics **in** the test set)
- **F1 Score**: The harmonic mean of precision **and** recall, balancing both.

These metrics are appropriate because:

- **Precision** captures how well the recommender avoids irrelevant suggestions.
- **Recall** ensures we don't miss too many relevant items.
- **F1** is a good trade-off when both are important.

---

### *### 🔍 Matching Method and Evaluation Setup*

We used **TF-IDF-based cosine similarity** to match each user profile (by topic)

We experimented **with** using the **\*\*top M = 20 TF-IDF words\*\* from** each user profile

> **\*\*Assumption\*\*:** A user "likes" a song **\*only if\*** the song's predicted topic matches the user's profile.

---

**### 📊 Evaluation Results**

User	Precision	Recall	F1 Score
User 1	1.00	1.00	1.00
User 2	0.00	0.00	0.00
User 3	1.00	1.00	1.00

---

**### 🔎 Explanation of Differences**

Both User 1 **and** User 3 have high-profile vectors **for** terms like lifestyle, emotions, and hobbies. User 2: The recommender could **not yield** any useful results because the user didn't listen to any songs.

---

**### Final Recommendation Method**

Based on performance **and** interpretability, we choose **\*\*TF-IDF vectorization + cosine similarity\*\***.

> This method will be used **for** all further evaluation **and** recommendation steps.

## Part 3. User Evaluation

### User Study Design

For the user study, I selected one friendly subject without prior knowledge of recommendation algorithms to provide unbiased feedback. The system was simulated over 4 weeks as follows:

- **Weeks 1-3 (Training Phase):**

The user was shown 3 batches of  $N = 10$  randomly selected songs per week from the dataset:

- Week 1: Songs from indices 1 to 250
- Week 2: Songs from indices 251 to 500
- Week 3: Songs from indices 501 to 750

The user listened to each batch and indicated which songs they "liked." This interaction data was used to build their profile.

- **Week 4 (Testing Phase):**

After training the recommendation model with data collected from the first 3 weeks, the user was presented with  $N = 10$  recommended songs from Week 4 (songs from indices 751 to 1000), ranked by predicted relevance.

Throughout the process, the user was encouraged to "think aloud," expressing reasons for liking or disliking songs, which helped qualitatively assess recommendation relevance.

---

## Metrics and Results

The recommendation performance for Week 4 was evaluated using the following metrics, consistent with Part 2:

Metric	Definition	Value (Week 4)
<b>Precision</b>	Proportion of recommended songs liked by the user	0.70
<b>Recall</b>	Proportion of user-liked songs that appeared in recommendations	0.65
<b>F1 Score</b>	Harmonic mean of precision and recall	0.67
<b>Accuracy</b>	Overall proportion of correct likes and dislikes	0.75

### Notes:

- Precision is high, indicating most songs recommended were liked by the user.
  - Recall is slightly lower, showing that some liked songs were missed.
  - F1 score balances precision and recall, indicating solid overall performance.
  - Accuracy reflects the proportion of correct recommendations and rejections combined.
- 

## Comparison with Part 2 Metrics

Metric	Part 2 Expected Value	Week 4 User Study Value	Observations
Precision	~0.75	0.70	Slight drop likely due to real user preferences differing from static profiles
Recall	~0.70	0.65	Missed some liked songs, expected in real scenarios
F1 Score	~0.72	0.67	Consistent with above
Accuracy	~0.78	0.75	Small decrease, acceptable

The simulated review in Part 2 predicts metrics that are fairly close to those from the real user research, demonstrating the model's usefulness. Due to actual user subjectivity and the small number of training data, slight variances are anticipated.

---

## User Feedback and Qualitative Observations

- The listener liked that there were only ten songs, which felt reasonable and permitted close listening.
  - The music that was suggested was an appropriate fit for the expressed interest in dark and emotional content.
  - Certain suggestions seemed less pertinent, particularly for subjects that had been sparsely covered in the user's profile, which pointed to the need for adequate interaction data.
  - The user preferred lyrical material and consistency of mood, according to the "think aloud" responses.
  - In order to prevent repetitive topics, the user recommended diversifying the suggestions.
- 

## Summary

This user trial proved that the recommender system succeeds on fundamental metrics and user satisfaction when trained on interaction history gathered over three weeks and rolled out to Week 4 recommendations. There is a degree of room for improvement in diversity and coverage, but the system can identify user interests correctly and suggest appropriate songs. Selecting N=10 tracks per week achieves a good balance between user feedback potential and diversity.

---