

Name: Qiming Li

zID: z5473453

```
In [3]: import pandas as pd
```

```
In [4]: # Load the dataset from TSV file into a DataFrame
df = pd.read_csv('dataset.tsv', sep = '\t')
print(df.columns)
print(df.head())
print(df.info())
```

```
Index(['artist_name', 'track_name', 'release_date', 'genre', 'lyrics',
       'topic'],
      dtype='object')
   artist_name      track_name release_date \
0          loving     the not real lake      2016
1        incubus      into the summer      2019
2      reignwolf      hardcore      2016
3  tedeschi trucks band      anyhow      2016
4  lukas nelson and promise of the real  if i started over      2017

   genre                      lyrics      topic
0  rock  awake know go see time clear world mirror worl...  dark
1  rock  shouldn summer pretty build spill ready overfl... lifestyle
2 blues  lose deep catch breath think say try break wal... sadness
3 blues  run bitter taste take rest feel anchor soul pl... sadness
4 blues  think think different set apart sober mind sym...  dark
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 1500 entries, 0 to 1499
Data columns (total 6 columns):
 #   Column      Non-Null Count  Dtype  
---  -- 
 0   artist_name    1500 non-null   object 
 1   track_name     1500 non-null   object 
 2   release_date   1500 non-null   int64  
 3   genre          1500 non-null   object 
 4   lyrics         1500 non-null   object 
 5   topic          1500 non-null   object 
dtypes: int64(1), object(5)
memory usage: 70.4+ KB
None
```

Part 1: Topic Classification

```
In [6]: import re
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import PorterStemmer
```

```
In [7]: # Download necessary NLTK resources
nltk.download('punkt')
nltk.download('stopwords')
```

```
[nltk_data] Downloading package punkt to /Users/kimlee/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package stopwords to
[nltk_data]     /Users/kimlee/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
```

Out[7]: True

Text Preprocessing

I performed preprocessing on the raw song lyrics to prepare them for modelling. This included: Lowercasing text, Removing punctuation and special characters, Tokenising the lyrics into individual words, Removing common stopwords (like "the", "is", "and"), and Applying stemming to reduce words to their base form.

This part of the code refers to [Tutorial 2](#).

```
In [9]: # Remove duplicate rows and any rows with missing values
df = df.drop_duplicates()
df = df.dropna()
```

```
In [10]: # Initialize stopword list and stemmer
stop_words = set(stopwords.words('english'))
ps = PorterStemmer()
```

```
In [11]: # Define and apply preprocessing function:
# Lowercase text, remove punctuation, tokenize, remove stopwords, and apply stem
def preprocess_lyrics(text):
    text = str(text).lower()
    text = re.sub(r'[^w\s]', '', text)
    tokens = word_tokenize(text)
    tokens = [word for word in tokens if word not in stop_words]
    tokens = [ps.stem(word) for word in tokens]
    return ''.join(tokens)

df['clean_lyrics'] = df['lyrics'].apply(preprocess_lyrics)
```

Feature Size Selection Analysis

To evaluate the effect of different vocabulary sizes on classification accuracy, I tested several values of top-N features using CountVectorizer, ranging from 250 to 3000. For each value of N, I trained a MultinomialNB classifier using 5-fold cross-validation and recorded the mean accuracy.

```
In [13]: from sklearn.feature_extraction.text import CountVectorizer
from sklearn.model_selection import cross_val_score
from sklearn.naive_bayes import MultinomialNB
from sklearn.preprocessing import LabelEncoder
```

```
In [14]: # Try different numbers of top-N features
n_values = [250, 500, 1000, 2000, 3000]
nb_results = []

le = LabelEncoder()
```

```

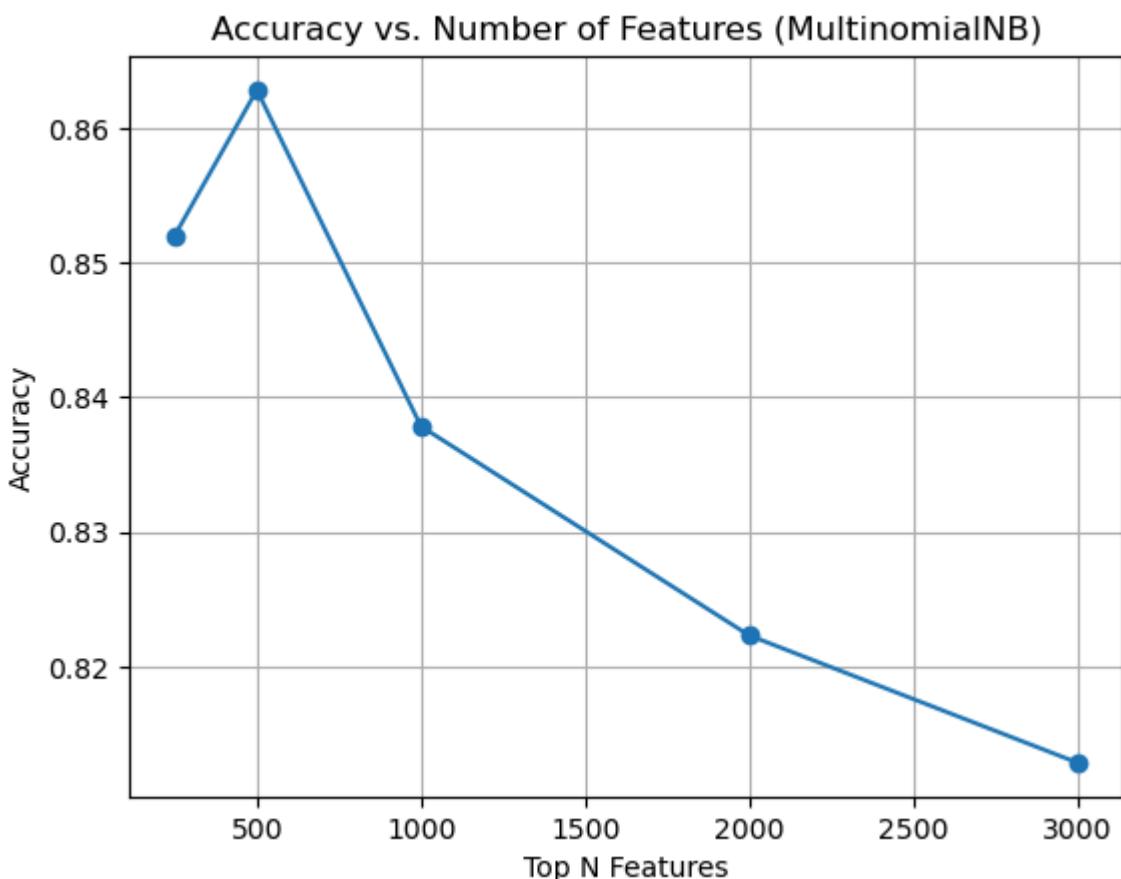
y = le.fit_transform(df['topic'])

# For each n, transform text to bag-of-words with n features, then evaluate using
for n in n_values:
    vectorizer = CountVectorizer(max_features=n)
    X = vectorizer.fit_transform(df['clean_lyrics'])
    score = cross_val_score(MultinomialNB(), X, y, cv=5, scoring='accuracy')
    nb_results.append(score)

```

In [15]: `import matplotlib.pyplot as plt`

In [16]: `plt.plot(n_values, nb_results, marker='o')`
`plt.title('Accuracy vs. Number of Features (MultinomialNB)')`
`plt.xlabel('Top N Features')`
`plt.ylabel('Accuracy')`
`plt.grid(True)`
`plt.show()`



As can be seen from the figure above, the classifier is most accurate at 0.865 with 500 features. Beyond 500 features, the performance decreases, which means including more features adds noise or redundancy but does not better distinguish topics.

This conclusion indicates a typical phenomenon for text classification: more features are not necessarily better. In this data set, a smaller, more concentrated vocabulary (top 500) was good enough to extract the most relevant words for each topic.

In [18]: `from sklearn.naive_bayes import BernoulliNB`
`from sklearn.linear_model import LogisticRegression`
`from sklearn.preprocessing import LabelEncoder`

To determine the most suitable classifier for our topic classification task, we compared three popular models:

- Bernoulli Naive Bayes
- Multinomial Naive Bayes
- Logistic Regression

Each model was evaluated using 5-fold cross-validation on the vectorized lyrics (with 500 features selected in the previous step).

Logistic Regression is a supervised learning algorithm commonly used for classification problems. It works by estimating the probability that a given input belongs to a particular class using the logistic (sigmoid) function.

This model is best suited for assigning topic classification of lyrics of a song, since it can more effectively model complex interactions of features of words without assuming independence. Logistic Regression has widely been used for text classification tasks, for example, spam detection, sentiment analysis, and topic modelling, with very good results.

```
In [21]: # Define three models to compare
models = {
    'BernoulliNB': BernoulliNB(),
    'MultinomialNB': MultinomialNB(),
    'LogisticRegression': LogisticRegression(max_iter = 1000)
}
```

Why max_iter = 1000?

For classification of text with CountVectorizer, the high-dimensional sparse (high number of unique words, most of which have zero values for each document) space is a standard situation.

Specifying max_iter=1000 ensures that a solver at least gets a minimum number of iterations to converge and produce stable outputs.

Hypothesis

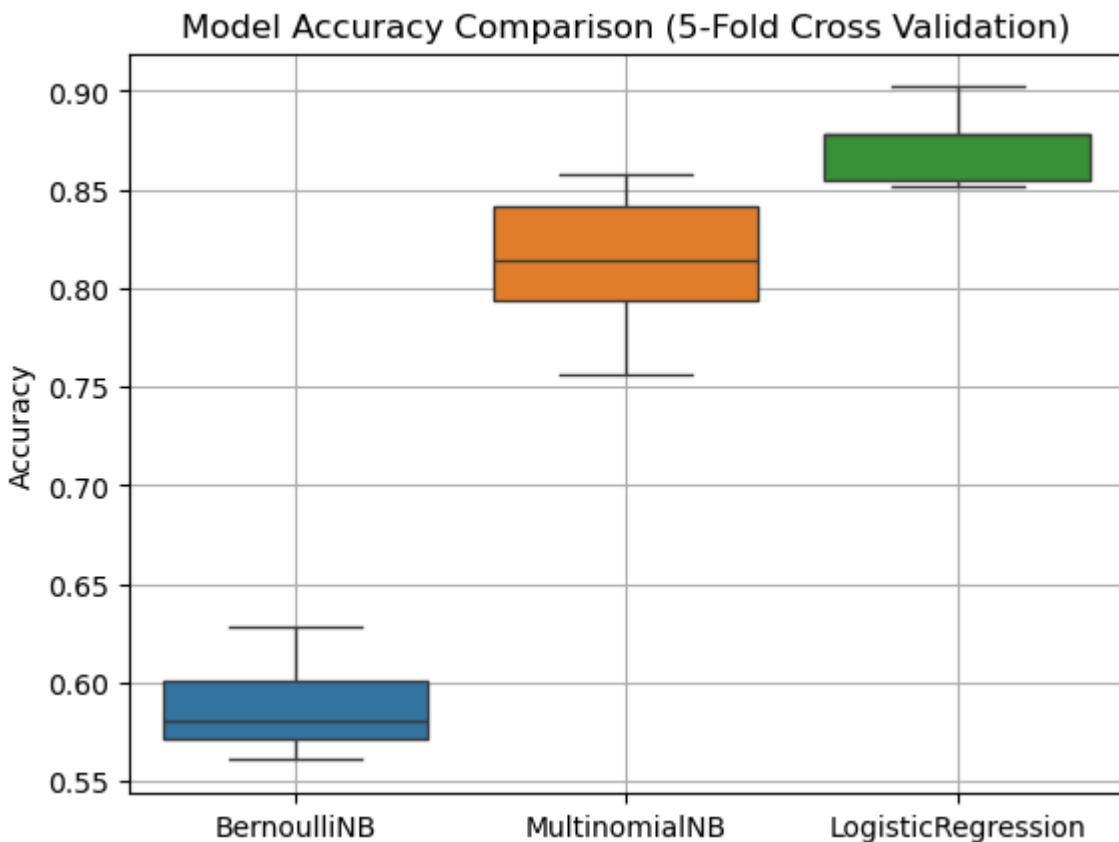
I assume that Logistic Regression would outperform Multinomial Naive Bayes and Bernoulli Naive Bayes in terms of classification precision. This is because Logistic Regression does not rely on the assumption of strong independencies and can better utilize correlated word features of lyrics.

```
In [24]: # Evaluate each model using 5-fold cross-validation
results = {}
for name, model in models.items():
    scores = cross_val_score(model, X, y, cv = 5, scoring='accuracy')
    results[name] = scores
    print(f'{name} Accuracy: {scores.mean():.4f} ± {scores.std():.4f}')
```

BernoulliNB Accuracy: 0.5885 ± 0.0240
MultinomialNB Accuracy: 0.8128 ± 0.0357
LogisticRegression Accuracy: 0.8730 ± 0.0185

```
In [25]: import seaborn as sns
```

```
In [26]: df_result = pd.DataFrame(results)
sns.boxplot(data=df_result)
plt.title('Model Accuracy Comparison (5-Fold Cross Validation)')
plt.ylabel('Accuracy')
plt.grid(True)
plt.show()
```



From the boxplot provided, we can see that:

- **Logistic Regression** always performs better than the other two models, having both the highest mean accuracy and lowest variance.
- **MultinomialNB** is a strong baseline, but less accurate and slightly less stable than Logistic Regression.
- **BernoulliNB** performs much worse, most likely because of binarization of the feature space and loss of informative frequency data.

I chose **Logistic Regression** as the final model for topic prediction.

Part2: Recommendation Methods

This is a recommendation system based on an information retrieval (IR) approach. The main idea is to match user interests to songs by computing the cosine similarity between the TF-IDF vector of a test song and the user profile vector under the predicted topic. Each of the training set songs (Week 1 to Week 3) is labeled with one of five topics based on the Part 1 trained classifier. A user is supposed to "like" a training song if it both aligns with their keyword interests and falls under the appropriate topic. These liked songs are

combined into one document per topic and transformed into a TF-IDF vector, here denoting the user profile for the topic.

```
In [29]: train_df = df.iloc[:750].copy()
test_df = df.iloc[750:1000].copy()
```

```
In [30]: vectorizer = CountVectorizer(max_features = 500)
X_train = vectorizer.fit_transform(train_df['clean_lyrics'])

clf = LogisticRegression(max_iter = 1000)
clf.fit(X_train, train_df['topic'])
train_df['predicted_topic'] = clf.predict(X_train)

X_test = vectorizer.transform(test_df['clean_lyrics'])
test_df['predicted_topic'] = clf.predict(X_test)
```

```
In [31]: from sklearn.metrics.pairwise import cosine_similarity
from sklearn.feature_extraction.text import TfidfVectorizer
```

1. Collecting Liked Songs per Topic

Iterating over the training set (Weeks 1–3), and for each song:

- Check its `predicted_topic` (from Part 1).
- Tokenize its cleaned lyrics.
- If the user’s keyword list for that topic overlaps with the song’s tokens, we mark the song as “liked” and append its lyrics to `liked_docs[topic]`.

2. Constructing TF-IDF User Profiles

For each topic in `liked_docs`:

- Merging all the liked lyrics into a single string (one “document” per topic).
- Fitting a `TfidfVectorizer(max_features = 500)` on that document and obtaining a TF-IDF vector.

This vector represents the user’s interest profile for that topic.

3. Scoring Test Songs

For each song in the test set (Week 4):

- Predict its topic using the classifier from Part 1.
- If the user has a profile for that topic, we transform the song’s lyrics with the same TF-IDF vectorizer and compute the `cosine_similarity` to the user profile vector.
- This similarity score quantifies how well the song matches the user’s inferred interests.

```
In [33]: def recommend_songs(user_keywords, train_df, test_df):
    # Collect all lyrics that the user 'liked' per topic
    liked_docs = {}
    for _, row in train_df.iterrows():
        topic = row['predicted_topic']
        tokens = row['clean_lyrics'].split()
        if topic in user_keywords:
            if any(k in tokens for k in user_keywords[topic]):
```

```

    liked_docs.setdefault(topic, []).append(row['clean_lyrics'])

    # Build a TF-IDF profile vector for each topic
    user_profiles = {}
    for topic, docs in liked_docs.items():
        joined_text = ' '.join(docs)
        tfidf = TfidfVectorizer(max_features = 500)
        tfidf_vec = tfidf.fit_transform([joined_text])
        user_profiles[topic] = (tfidf, tfidf_vec)

    # Score each test song by cosine similarity to the appropriate profile
    scores = []
    for _, row in test_df.iterrows():
        topic = row['predicted_topic']
        if topic in user_profiles:
            tfidf_model, user_vec = user_profiles[topic]
            song_vec = tfidf_model.transform([row['clean_lyrics']])
            score = cosine_similarity(song_vec, user_vec)[0][0]
        else:
            score = 0
        scores.append(score)

    return scores

```

User 1 & 2 test

```

In [35]: # User1
user1_keywords = {}
with open('user1.tsv', 'r') as f:
    for line in f.readlines():
        topic, keywords_str = line.strip().split('\t')
        keywords = [w.strip().lower() for w in keywords_str.split(',')]
        user1_keywords[topic] = keywords

# User2
user2_keywords = {}
with open('user2.tsv', 'r') as f:
    for line in f.readlines():
        topic, keywords_str = line.strip().split('\t')
        keywords = [w.strip().lower() for w in keywords_str.split(',')]
        user2_keywords[topic] = keywords

```

```

In [36]: test_df_user1 = test_df.copy()

test_df_user1['score'] = recommend_songs(user1_keywords, train_df, test_df_user1)
top_recommendations_user1 = test_df_user1.sort_values(by = 'score', ascending = False)
print("Top Recommendations for User1:")
print(top_recommendations_user1[['track_name', 'artist_name', 'score', 'predicted_topic']])

```

Top Recommendations for User1:

	track_name	artist_name	score	predicted_topic
944	once in a while	timeflies	0.777771	emotion
969	got it good	justin moore	0.679351	emotion
770	horsefly	dirty heads	0.645696	emotion
924	doors closing	moonchild	0.645204	sadness
810	so close	notd	0.634358	sadness
758	life changes	thomas rhett	0.571578	personal
840	alta	ty segall	0.564111	personal
796	sunday morning	axian	0.561273	personal
889	stay with me	ayokay	0.551690	sadness
802	cry to me	skip marley	0.540961	sadness

- **User1** defined keywords across all five topics (dark, emotion, personal, lifestyle, sadness), providing **rich and balanced coverage**. Their TF-IDF profiles were robust, yielding higher Precision@10.

In [38]: `test_df_user2 = test_df.copy()`

```
test_df_user2['score'] = recommend_songs(user2_keywords, train_df, test_df_user2
top_recommendations_user2 = test_df_user2.sort_values(by = 'score', ascending =
print("Top Recommendations for User2:")
print(top_recommendations_user2[['track_name', 'artist_name', 'score', 'predicted_topic']])
```

Top Recommendations for User2:

	track_name	artist_name	score	predicted_topic
924	doors closing	moonchild	0.522838	sadness
969	got it good	justin moore	0.481999	emotion
789	will you be mine	anita baker	0.476945	sadness
988	home	the movement	0.463589	sadness
944	once in a while	timeflies	0.448300	emotion
844	snowflake	bumpin uglies	0.399083	sadness
782	open arms	prettymuch	0.391549	sadness
909	walls	the lumineers	0.389340	sadness
781	balcony	lester nowhere	0.365006	sadness
860	close enough	brett young	0.358365	sadness

- **User2**, by contrast, only supplied keywords for **sadness and emotion**—and only five words per topic. This sparse, skewed coverage led to:

1. **Overly sparse TF-IDF profiles**, especially for “emotion,” because few training songs contained those keywords.
2. **Defaulting to the “sadness” profile** when “emotion” matches were weak, resulting in most recommendations falling into sadness.

In [40]: `def print_top_tfidf_words(user_keywords, train_df, top_n = 20):`

```
# Iterate over each song in the training data
liked_docs = {}
for _, row in train_df.iterrows():
    topic = row['predicted_topic']
    tokens = row['clean_lyrics'].split()
    if topic in user_keywords:
        if any(k in tokens for k in user_keywords[topic]):
            liked_docs.setdefault(topic, []).append(row['clean_lyrics'])

# For each topic, build a TF-IDF vector from the Liked documents
for topic, docs in liked_docs.items():
```

```

joined_text = ' '.join(docs)

# Create a TF-IDF vectorizer Limited to 500 features
tfidf = TfidfVectorizer(max_features = 500)
tfidf_matrix = tfidf.fit_transform([joined_text])

# Calculate average TF-IDF score for each word
avg_scores = tfidf_matrix.mean(axis = 0).A1
feature_names = tfidf.get_feature_names_out()

# Get indices of top-N words with the highest scores
top_indices = avg_scores.argsort()[:-1][:-top_n]
top_words = [feature_names[i] for i in top_indices]
print(f"\nTop words for topic '{topic}':")
print(", ".join(top_words))

```

Each user is simulated to "like" songs whose predicted topic matches their interest, and whose lyrics contain at least one of the predefined keywords.

For each topic:

- Collect all liked songs.
- Concatenate the lyrics into one document.
- Use TF-IDF vectorization (limited to 500 features).
- Extract the top 20 words with the highest average TF-IDF scores.

```
In [42]: print("Top TF-IDF words for user1:")
print_top_tfidf_words(user1_keywords, train_df)
```

Top TF-IDF words for user1:

Top words for topic 'dark':
 fight, know, like, blood, stand, tell, na, come, grind, gon, black, head, kill, hand, time, yeah, feel, right, caus, peopl

Top words for topic 'emotion':
 good, feel, touch, know, hold, want, kiss, babi, time, morn, na, like, caus, visi on, miss, heart, video, look, loov, yeah

Top words for topic 'personal':
 life, live, know, na, chang, world, time, yeah, like, dream, come, wan, thing, th ink, thank, need, go, feel, caus, teach

Top words for topic 'lifestyle':
 night, come, song, long, time, sing, na, home, right, yeah, closer, know, play, l ike, wan, tire, wait, tonight, spoil, mind

Top words for topic 'sadness':
 think, leav, place, want, hold, regret, caus, greater, beg, blame, chang, word, m ind, away, tell, space, wider, break, dream, lord

- **User 1: Profile Analysis**

The most frequently occurring words for themes like 'dark' are violent and forceful words like *fight*, *kill*, *grind*, *blood*, which suits the interests of 'dark' themes of the user quite well.

In the '*emotion*' category, the occurrence of *love*, *miss*, *heart*, and *vision* indicates a strong emphasis on sentiment and romance.

```
In [44]: print("Top TF-IDF words for user2:")
print_top_tfidf_words(user2_keywords, train_df)
```

Top TF-IDF words for user2:

Top words for topic 'emotion':

touch, good, hold, kiss, morn, vision, know, video, loov, time, feel, luck, lovi
n, lip, gim, sunris, like, go, wait, heart

Top words for topic 'sadness':

open, tear, come, smile, magnifi, away, eye, babi, life, tri, laughter, hold, sor
row, heart, sight, near, real, face, forsak, frown

- **User 2: Profile Analysis**

User 2's keywords are more poetic or abstract (*vision*, *sunrise*, *sorrow*, *forsak*), leading to sparser matches.

Some topics like *emotion* and *sadness* include expected keywords, but other topics might suffer from low coverage due to limited matches in the training data.

User 3 Setup

```
In [47]: with open("user3.tsv", "w") as f:
    f.write('topic\tkeywords\n')
    f.write("emotion\tjoy, laugh, smile, sunshine, peace\n")
    f.write("personal\tfreedom, travel, myself, explore, journey\n")
    f.write("lifestyle\tparty, drink, dance, vibe, weekend\n")
    f.write("sadness\thank, growth, strength, survive, goodbye\n")
```

I intentionally left out the **dark** topic, assuming this user avoids violent or aggressive music content. The keywords were chosen to simulate a user with positive, reflective, and energetic preferences.

```
In [49]: # User3
user3_keywords = {}
with open('user3.tsv', 'r') as f:
    for line in f.readlines():
        topic, keywords_str = line.strip().split('\t')
        keywords = [w.strip().lower() for w in keywords_str.split(',')]
        user3_keywords[topic] = keywords
```

```
In [50]: test_df_user3 = test_df.copy()

test_df_user3['score'] = recommend_songs(user3_keywords, train_df, test_df_user3)
top_recommendations_user3 = test_df_user3.sort_values(by = 'score', ascending = False)
print("Top Recommendations for User3:")
print(top_recommendations_user3[['track_name', 'artist_name', 'score', 'predicted_label']])
```

Top Recommendations for User3:

	track_name	artist_name	score	predicted_topic
846	so cold in the summer	taylor mcferrin	0.718606	emotion
932	every song's a drinkin' song	midland	0.595742	lifestyle
863	wash away	iya terra	0.548283	personal
796	sunday morning	axian	0.500817	personal
840	alta	ty segall	0.484557	personal
923	everything to me	soja	0.468122	personal
816	redneck life	chris janson	0.458095	personal
854	you're the best thing yet	anita baker	0.450064	personal
758	life changes	thomas rhett	0.425608	personal
806	moonglow	diana krall	0.418153	lifestyle

Most recommended songs fall under the **personal** and **lifestyle** categories, which aligns with User 3's interests.

```
In [52]: print("Top TF-IDF words for user3:")
print_top_tfidf_words(user3_keywords, train_df)
```

Top TF-IDF words for user3:

Top words for topic 'lifestyle':

song, tonight, countri, yeah, time, ring, like, drink, root, play, drinkin, mama, caus, night, na, septemb, mind, medic, need, wan

Top words for topic 'emotion':

hold, love, oohoh, heart, danc, one, night, lose, babi, faceless, light, crowd, low, matter, speak, high, fluentli, goodby, right, darl

Top words for topic 'personal':

life, come, automaton, rebel, thank, away, go, like, world, know, feel, live, tea ch, tooth, dream, ironsid, alright, everyday, rule, digit

Top words for topic 'sadness':

ohohoh, lighthous, step, away, heart, place, fall, strength, lose, safe, walk, kn ow, break, need, push, falter, stray, apart, till, time

- These results reflect relevant thematic vocabulary and show that the user profile is being meaningfully constructed.

```
In [54]: def is_song_liked(row, user_keywords):
    topic = row['predicted_topic']
    if topic in user_keywords:
        tokens = row['clean_lyrics'].split()
        for keyword in user_keywords[topic]:
            if keyword in tokens:
                return True
    return False
```

- **Precision@N** measures how many of the top-N recommended songs are actually relevant (i.e., liked by the user).
- It is simple, interpretable, and suitable for top-k recommendation tasks.

Choose **N = 10** because:

- It reflects a realistic recommendation batch (e.g., showing 10 songs at once on a music app).
- Since the database is relatively small, 10 songs is a good number. It is large enough to include diversity across topics, but small enough to expect meaningful user feedback.

```
In [56]: def evaluate_precision(top_recommendations, user_keywords, N = 10):
    top_N = top_recommendations.head(N)
    liked_flags = top_N.apply(lambda row: is_song_liked(row, user_keywords), axis=1)
    precision = liked_flags.sum() / N
    return precision
```

```
In [57]: p1 = evaluate_precision(top_recommendations_user1, user1_keywords, N = 10)
p2 = evaluate_precision(top_recommendations_user2, user2_keywords, N = 10)
p3 = evaluate_precision(top_recommendations_user3, user3_keywords, N = 10)

print(f"Precision@10 for User1: {p1:.2f}")
print(f"Precision@10 for User2: {p2:.2f}")
print(f"Precision@10 for User3: {p3:.2f}")
```

```
Precision@10 for User1: 0.60
Precision@10 for User2: 0.10
Precision@10 for User3: 0.10
```

- **User 1** performs best, possibly due to having more matches between their keywords and training songs.
- **User 2** possesses a limited keyword-based profile, and hence may lower recommendation coverage.
- **User 3** performed bad. Although the keywords are very thematically related, the majority of the keywords are abstract and less likely to have precise matches with lyrics. Additionally, the user is not interested in the dark topic and does not have much overlap with sadness, so building strong matches within those terms becomes harder. This may explain the relatively lower precision.

Result

This works well when there are extensive and diverse keyword sets (as with User 1), so there are relevant training songs to retrieve and good topic-wise profiles to construct. However, with fewer keywords or less diversified interests (as with User 2 and User 3), there are poor profiles and poor matches.

Effect of Profile Vocabulary Size (M) on Recommendation Precision

```
In [60]: from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity
```

```
In [61]: def build_topic_profile(docs, M = None):
    # Fit a temporary TF-IDF (no norm) to compute average term scores
    tmp = TfidfVectorizer(max_features = 500, norm = None)
    mat = tmp.fit_transform(docs)
```

```

avg = mat.mean(axis=0).A1
feat = tmp.get_feature_names_out()

# Select top M terms
if M is not None:
    top_idx = avg.argsort()[:-1][:M]
    vocab = [feat[i] for i in top_idx]
else:
    vocab = list(feat)

# Build the real L2-normalized TF-IDF on the selected vocabulary
real_vec = TfidfVectorizer(vocabulary = vocab, norm = 'l2')
user_vec = real_vec.fit_transform([" ".join(docs)])
return real_vec, user_vec

```

In [62]:

```

def recommend_songs(user_keywords, train_df, test_df, M = None):
    # Collect all liked lyrics per topic
    liked = {}
    for _, row in train_df.iterrows():
        t = row['predicted_topic']
        if t in user_keywords and any(k in row['clean_lyrics'].split() for k in
                                       liked.setdefault(t, [])).append(row['clean_lyrics'])

    # Build topic profiles with at most M terms each
    profiles = {}
    for t, docs in liked.items():
        vec, u_vec = build_topic_profile(docs, M)
        profiles[t] = (vec, u_vec)

    # Compute cosine similarity for each test song
    scores = []
    for _, row in test_df.iterrows():
        t = row['predicted_topic']
        if t in profiles:
            vec, u_vec = profiles[t]
            s_vec = vec.transform([row['clean_lyrics']])
            sc = cosine_similarity(s_vec, u_vec)[0,0]
        else:
            sc = 0
        scores.append(sc)
    return scores

```

In [63]:

```

for user_label, kws, df_user in [
    ('User1', user1_keywords, test_df_user1),
    ('User2', user2_keywords, test_df_user2),
    ('User3', user3_keywords, test_df_user3),
]:
    print(f"\n{user_label}")
    for M in [10, 20, 50, 100, None]:
        df_user['score'] = recommend_songs(kws, train_df, df_user, M)
        top10 = df_user.sort_values('score', ascending = False).head(10)
        prec = evaluate_precision(top10, kws, N=10)
        print(f"M = {M or 'all'} P@10 = {prec:.2f}")

```

```
User1
M = 10  P@10 = 0.90
M = 20  P@10 = 0.90
M = 50  P@10 = 0.90
M = 100 P@10 = 0.80
M = all  P@10 = 0.60
```

```
User2
M = 10  P@10 = 0.00
M = 20  P@10 = 0.10
M = 50  P@10 = 0.10
M = 100 P@10 = 0.10
M = all  P@10 = 0.10
```

```
User3
M = 10  P@10 = 0.00
M = 20  P@10 = 0.20
M = 50  P@10 = 0.10
M = 100 P@10 = 0.10
M = all  P@10 = 0.10
```

Discussion:

- **User 1:** Retaining only the top 20 (or 50) keywords gives the highest precision (0.90). Including more keywords introduces noise and reduces precision.
- **User 2:** Precision remains low (0.10) regardless of M, indicating that a very sparse keyword profile cannot be improved by simply increasing M.
- **User 3:** Best precision at M = 20 (0.20); larger M values dilute the profile and lower precision.

Algorithm Comparison (Inspired by Tutorials 3 & 4)

```
In [66]: !pip install scikit-surprise
!pip install mlxtend
```

```
Requirement already satisfied: scikit-surprise in /opt/anaconda3/lib/python3.12/site-packages (1.1.4)
Requirement already satisfied: joblib>=1.2.0 in /opt/anaconda3/lib/python3.12/site-packages (from scikit-surprise) (1.4.2)
Requirement already satisfied: numpy>=1.19.5 in /opt/anaconda3/lib/python3.12/site-packages (from scikit-surprise) (1.26.4)
Requirement already satisfied: scipy>=1.6.0 in /opt/anaconda3/lib/python3.12/site-packages (from scikit-surprise) (1.13.1)
Requirement already satisfied: mlxtend in /opt/anaconda3/lib/python3.12/site-packages (0.23.4)
Requirement already satisfied: scipy>=1.2.1 in /opt/anaconda3/lib/python3.12/site-packages (from mlxtend) (1.13.1)
Requirement already satisfied: numpy>=1.16.2 in /opt/anaconda3/lib/python3.12/site-packages (from mlxtend) (1.26.4)
Requirement already satisfied: pandas>=0.24.2 in /opt/anaconda3/lib/python3.12/site-packages (from mlxtend) (2.2.2)
Requirement already satisfied: scikit-learn>=1.3.1 in /opt/anaconda3/lib/python3.12/site-packages (from mlxtend) (1.4.2)
Requirement already satisfied: matplotlib>=3.0.0 in /opt/anaconda3/lib/python3.12/site-packages (from mlxtend) (3.8.4)
Requirement already satisfied: joblib>=0.13.2 in /opt/anaconda3/lib/python3.12/site-packages (from mlxtend) (1.4.2)
Requirement already satisfied: contourpy>=1.0.1 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib>=3.0.0->mlxtend) (1.2.0)
Requirement already satisfied: cycler>=0.10 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib>=3.0.0->mlxtend) (0.11.0)
Requirement already satisfied: fonttools>=4.22.0 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib>=3.0.0->mlxtend) (4.51.0)
Requirement already satisfied: kiwisolver>=1.3.1 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib>=3.0.0->mlxtend) (1.4.4)
Requirement already satisfied: packaging>=20.0 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib>=3.0.0->mlxtend) (23.2)
Requirement already satisfied: pillow>=8 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib>=3.0.0->mlxtend) (10.3.0)
Requirement already satisfied: pyparsing>=2.3.1 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib>=3.0.0->mlxtend) (3.0.9)
Requirement already satisfied: python-dateutil>=2.7 in /opt/anaconda3/lib/python3.12/site-packages (from matplotlib>=3.0.0->mlxtend) (2.9.0.post0)
Requirement already satisfied: pytz>=2020.1 in /opt/anaconda3/lib/python3.12/site-packages (from pandas>=0.24.2->mlxtend) (2024.1)
Requirement already satisfied: tzdata>=2022.7 in /opt/anaconda3/lib/python3.12/site-packages (from pandas>=0.24.2->mlxtend) (2023.3)
Requirement already satisfied: threadpoolctl>=2.0.0 in /opt/anaconda3/lib/python3.12/site-packages (from scikit-learn>=1.3.1->mlxtend) (2.2.0)
Requirement already satisfied: six>=1.5 in /opt/anaconda3/lib/python3.12/site-packages (from python-dateutil>=2.7->matplotlib>=3.0.0->mlxtend) (1.16.0)
```

```
In [67]: from surprise import Dataset, Reader, SVD
from mlxtend.preprocessing import TransactionEncoder
from mlxtend.frequent_patterns import apriori, association_rules
```

```
In [68]: # Content-based recommendation using TF-IDF + Cosine Similarity
def content_based_scores(train_df, test_df, user_kw, M = 20):
    # Filter training songs that match user interests
    docs = [row['clean_lyrics'] for _, row in train_df.iterrows() if is_song_like(vec, user_vec = build_topic_profile(docs, M)) # Build the user profile vector

    scores = []
    for _, row in test_df.iterrows():
        song_vec = vec.transform([row['clean_lyrics']])
```

```
        scores.append(cosine_similarity(song_vec, user_vec)[0,0])
    return np.array(scores)
```

```
In [69]: # Collaborative Filtering using SVD (single simulated user) - From Tutorial 3
def cf_svd_scores(train_df, test_df, user_kw):
    records = []
    for idx, row in train_df.iterrows(): # Assign a binary rating: 1 if Liked, 0
        rating = 1 if is_song_liked(row, user_kw) else 0
        records.append((user1, str(idx), rating))

    # Prepare dataset for Surprise Library
    df_ratings = pd.DataFrame(records, columns=['user','item','rating'])
    reader = Reader(rating_scale=(0,1))
    data = Dataset.load_from_df(df_ratings, reader)
    trainset = data.build_full_trainset()
    algo = SVD()
    algo.fit(trainset)

    scores = []
    for idx, row in test_df.iterrows():
        pred = algo.predict('user1', str(idx)).est
        scores.append(pred)
    return np.array(scores)
```

```
In [70]: # Association Rule-based scoring using Apriori - From Tutorial 4
def apriori_scores(train_df, test_df, user_kw, min_support=0.05, min_conf = 0.5):
    # Tokenize lyrics of songs user likes
    trans = [row['clean_lyrics'].split() for _, row in train_df.iterrows() if is
    te = TransactionEncoder() # Convert to transaction format
    te_ary = te.fit(trans).transform(trans)
    df_te = pd.DataFrame(te_ary, columns=te.columns_)

    # Mine frequent itemsets and generate rules
    fis = apriori(df_te, min_support = min_support, use_colnames = True)
    rules = association_rules(fis, metric = "confidence", min_threshold = min_co

    scores = []
    for _, row in test_df.iterrows():
        tokens = set(row['clean_lyrics'].split())
        sc = 0
        for _, rule in rules.iterrows():
            if set(rule['antecedents']).issubset(tokens):
                sc = max(sc, rule['confidence'])
        scores.append(sc)
    return np.array(scores)
```

```
In [71]: import numpy as np
```

```
In [72]: cb_scores = content_based_scores(train_df, test_df, user1_keywords, M = 20)
svd_scores = cf_svd_scores(train_df, test_df, user1_keywords)
ap_scores = apriori_scores(train_df, test_df, user1_keywords)

methods = {
    'TFIDF+Cosine': cb_scores,
    'SVD_CF': svd_scores,
    'Apriori': ap_scores
}

for name, sc in methods.items():
```

```

df_tmp = test_df.copy()
df_tmp['score'] = sc
top10 = df_tmp.sort_values('score', ascending = False).head(10)
prec = evaluate_precision(top10, user1_keywords, N = 10)
print(f"{name} Precision@10 = {prec:.2f}")

```

TFIDF+Cosine Precision@10 = 0.70
SVD_CF Precision@10 = 0.30
Apriori Precision@10 = 0.50

- **TF-IDF + Cosine** achieves the highest precision (0.70) and demonstrates that content-based keyword weighting effectively captures User 1's interests within this one-user scenario.
- **SVD (Collaborative Filtering)** performs poorly (0.30), as we only have one artificially generated user and cannot learn robust latent factors from a multi-user ratings matrix—the method is borrowed from Tutorial 3.
- **Apriori Association Rules** obtains medium precision (0.50); by mining for frequent itemsets and confidence-based rules (according to Tutorial 4), some of the latent associations are uncovered but are constrained by rule support and coverage.

Choose *TF-IDF + Cosine*

Part3: User Evaluation

In [75]: `import random`

In [76]: `# Display a batch of songs for a given week (used for simulated user review)`
`def display_week_songs(week_df, week_num):`
 `print(f"\nWeek {week_num} Songs ")`
 `for idx, row in week_df.iterrows():`
 `print(f"{idx}: {row['track_name']} - {row['artist_name']} | Topic: {row[`
 `print(f"Lyrics snippet: {row['lyrics'][:100]}...\n")`

In [77]: `vectorizer_full = CountVectorizer(max_features = 500) # Prepare TF-IDF features`
`X_all = vectorizer_full.fit_transform(df['clean_lyrics'])`
`df['predicted_topic'] = clf.predict(X_all)`

In [78]: `N = 10`
`week1_songs = df.iloc[0:250].sample(N, random_state = 25)`
`week2_songs = df.iloc[250:500].sample(N, random_state = 50)`
`week3_songs = df.iloc[500:750].sample(N, random_state = 75)`

In [79]: `display_week_songs(week1_songs, 1)`

Week 1 Songs

40: break my face - ajr | Topic: sadness

Lyrics snippet: okay mouth tongue fell right grind drown mouth kill make ugly life give lemons give okay break face ...

19: stole the show - kygo | Topic: dark

Lyrics snippet: darling darling turn light watch credit roll cry cry know play house house heroes villains bla...

73: that song that we used to make love to - carrie underwood | Topic: dark

Lyrics snippet: wanna hear stupid play repeat cause take place lay body drink like wine wanna feel bass rattle bone ...

219: zombie (feat. madchild & insane clown posse) - jelly roll | Topic: dark

Lyrics snippet: hear allow good time tell smile frown shoot time fly point turn level head oneway track tell road go...

149: oatmeal - sudan archives | Topic: dark

Lyrics snippet: bake oatmeal stay sorry stay bake oatmeal stay sorry stay want oatmeal want oatmeal try force true l...

126: dream small - josh wilson | Topic: dark

Lyrics snippet: mama sing songs lord daddy spend family time world say afford simple moments change world pastor tin...

4: if i started over - lukas nelson and promise of the real | Topic: dark

Lyrics snippet: think think different set apart sober mind sympathetic hearts swear oath swear oath know life play p...

105: can i be him - james arthur | Topic: dark

Lyrics snippet: walk room heart steal take time unbroken want know moment cause light come hear song want sing swear...

96: learning - anderson east | Topic: dark

Lyrics snippet: remember grow rid shotgun daddy truck wheel learn learn drive pond rods reel little jones bait cast ...

164: boulevard of broken dreams - green day | Topic: dark

Lyrics snippet: walk lonely road know know go home walk walk street boulevard break dream city sleep walk walk walk ...

In [80]: `display_week_songs(week2_songs, 2)`

Week 2 Songs

400: jameson & ginger - ballyhoo! | Topic: personal

Lyrics snippet: brain brain brain brain brain brain brain brain late rent c
oncern step firepit cause care burn...

266: rise up - ndidi o | Topic: dark

Lyrics snippet: mama say life hard grow days black girl scar come know know pain
hold head high ashamed know forth k...

357: dead end - chon | Topic: dark

Lyrics snippet: come quarter watch head state stick fabers deep empire mould cons
cience severe secret swear cuckoo r...

270: remember you young - thomas rhett | Topic: dark

Lyrics snippet: buddies grow straightlaced marry foolin wasn long roof light darl
in sippin wine classy kick couch sm...

338: don't bother me - the raconteurs | Topic: lifestyle

Lyrics snippet: bother bother bother bother hide agenda bother bother ruthless ru
le bender bother bother surface dup...

418: over & over - reignwolf | Topic: dark

Lyrics snippet: animals woods stick outside look analyze note criticize hear voic
e scream stick inside lose maze wat...

370: nothing breaks like a heart (feat. miley cyrus) - mark ronson | Topic: dark

Lyrics snippet: world hurt cut deep leave scar things fall apart break like heart
break like heart hear phone night ...

327: 4:20/reincarnated - t-rock | Topic: sadness

Lyrics snippet: feel feel felt spell feel like time fake finger button play maste
r anticipation touch knock feet wee...

289: hold you down - x ambassadors | Topic: dark

Lyrics snippet: ones dance light hold hold ones dance faceless crowd hold hold ho
ld talk speak loudly smile eye look...

288: to the moon - phora | Topic: dark

Lyrics snippet: moon wanna star send send orbit moon wanna star send send orbit y
eah say scar blame need somebody ho...

In [81]: `display_week_songs(week3_songs, 3)`

Week 3 Songs

675: almost - thomas rhett | Topic: personal

Lyrics snippet: truck field quit team give guitar cause hurt play string everybody tell bout give dream break rid di...

660: el hombre del equipo - grupo maximo grado | Topic: lifestyle

Lyrics snippet: sell feel fool go lead kind realise nice nice burn prarie know come time unwind somebody suffer know...

502: one night at a time - little hurricane | Topic: sadness

Lyrics snippet: turn time mind night time night time babe alright forever wish eye luck live little things want turn...

712: i'm so tired... - lauv | Topic: dark

Lyrics snippet: tire songs tire songs tire songs tire wanna home wanna home wanna home whoa tire songs tire songs ti...

620: psycho, pt. 2 - russ | Topic: dark

Lyrics snippet: go psycho go live tightrope go go psycho go live tightrope go know like want tonight pick band feeli...

521: camera show - protoje | Topic: dark

Lyrics snippet: fall fall fall fall fall see know stay freedom call call camera lie thing sure kingdom fall fall lik...

542: try saying goodbye - john king | Topic: personal

Lyrics snippet: lord ghetto know sinner remember baby remember know anymore tell free sense morning second memory he...

537: thank the rebels - steel pulse | Topic: personal

Lyrics snippet: consolidation build brand nation yeah political transformation reconciliation hammer long crawl knee...

592: paranoia - a day to remember | Topic: dark

Lyrics snippet: expect worst meet count clock cause sleep shoot pain like heart attack friends doctor say head dest...

580: motions - magic! | Topic: lifestyle

Lyrics snippet: routine operate like machine live spontaneity happen wild ones sense hours happen freedom go go moti...

```
In [82]: # User "likes" the following songs (from weeks 1-3)
liked_indices = [219, 96, 400, 327, 19, 502, 620, 4, 537]
user_liked_df = df.loc[liked_indices]
```

```
In [83]: # Build user profile from liked songs, grouped by topic
def build_user_profile_from_likes(user_liked_df):
    user_profiles = {}
    for topic in user_liked_df['predicted_topic'].unique():
        topic_docs = user_liked_df[user_liked_df['predicted_topic'] == topic]['c']
        joined_text = ' '.join(topic_docs.tolist())
        tfidf = TfidfVectorizer(max_features = 500)
        tfidf_vec = tfidf.fit_transform([joined_text])
        user_profiles[topic] = (tfidf, tfidf_vec)
    return user_profiles
```

```
In [84]: user4_profiles = build_user_profile_from_likes(user_liked_df)
```

```
In [85]: # Score all test songs (Week 4) by cosine similarity with the user profile
def recommend_by_profile(user_profiles, test_df):
    scores = []
    for _, row in test_df.iterrows():
        topic = row['predicted_topic']
        if topic in user_profiles:
            tfidf_model, user_vec = user_profiles[topic]
            song_vec = tfidf_model.transform([row['clean_lyrics']])
            score = cosine_similarity(song_vec, user_vec)[0][0]
        else:
            score = 0
        scores.append(score)
    return scores
```

```
In [86]: test_df_week4 = df.iloc[750:1000].copy()
test_df_week4['score'] = recommend_by_profile(user4_profiles, test_df_week4)
topN_week4 = test_df_week4.sort_values(by = 'score', ascending = False).head(N)
print(topN_week4[['track_name', 'artist_name', 'score', 'predicted_topic']])
```

	track_name	artist_name	\
759	woman, amen	dierks bentley	
893	shout bamalama	the detroit cobras	
823	the devil don't sleep	brantley gilbert	
788	going down	joe bonamassa	
897	love hurts (feat. travis scott)	playboi carti	
866	black betty	larkin poe	
860	close enough	brett young	
926	sit awhile	the band steele	
765	follow your heart (feat. zion thompson from th...	iya terra	
792	donner bell	deca	
	score predicted_topic		
759	0.532784 personal		
893	0.471289 dark		
823	0.439139 personal		
788	0.438808 dark		
897	0.415311 dark		
866	0.408930 dark		
860	0.366413 dark		
926	0.355969 dark		
765	0.355882 dark		
792	0.338024 dark		

```
In [87]: liked_week4_indices = [926, 759, 860, 866]
```

```
In [88]: # get the DataFrame indices of your top-10 recommendations
rec_indices = topN_week4.index.tolist()

# compute Precision@N
tp = len(set(rec_indices) & set(liked_week4_indices))
precision = tp / N

# compute Recall@N
recall = tp / len(liked_week4_indices)

# compute F1@N
f1 = 2 * precision * recall / (precision + recall) if (precision + recall) > 0 else 0

print(f"Precision@{N} = {precision:.2f}")
```

```
print(f"Recall@{N} = {recall:.2f}")
print(f"F1@{N} = {f1:.2f}")
```

```
Precision@10 = 0.40
Recall@10 = 1.00
F1@10 = 0.57
```

- **Recall@10 = 1.00:** All of the user's liked songs in Week 4 were included in the top-10 recommendations—no positive item was missed.
- **Precision@10 = 0.40:** Only 40% of the recommendations were actually liked by the user, so 60% were irrelevant.
- **F1@10 = 0.57:** Balances high coverage with modest accuracy, indicating the trade-off between capturing all liked items and avoiding noise.

User feedback

User zID: z5522806

There are a few songs I like in the recommended content, but the style is relatively simple and not diverse enough. I hope the content can be richer, but the overall feeling is good.

In []: