

```
In [ ]: ZID: z5508588  
Name: Ziyu Xiong  
COMP9727 Assignment1
```

```
In [132...]  
  
import pandas as pd  
  
# read dataset.tsv  
all_data = pd.read_csv('dataset.tsv', sep='\t')  
  
# View Data Format  
print(all_data.head())  
print(all_data.columns)
```

```
          artist_name      track_name release_date \
0             loving    the not real lake        2016
1            incubus   into the summer       2019
2           reignwolf     hardcore        2016
3      tedeschi trucks band      anyhow        2016
4  lukas nelson and promise of the real  if i started over       2017  
  
      genre          lyrics      topic
0  rock  awake know go see time clear world mirror worl...  dark
1  rock  shouldn summer pretty build spill ready overfl... lifestyle
2 blues  lose deep catch breath think say try break wal... sadness
3 blues  run bitter taste take rest feel anchor soul pl... sadness
4 blues  think think different set apart sober mind sym...  dark
Index(['artist_name', 'track_name', 'release_date', 'genre', 'lyrics',
       'topic'],
      dtype='object')
```

```
In [134...]  
  
# Combine into one text column  
def combine_fields(row):  
    return f'{row['artist_name']} {row['track_name']} {row['genre']} {row['lyric'}  
  
all_data['text'] = all_data.apply(combine_fields, axis=1)  
print(all_data[['text', 'topic']].head(3))
```

```
          text      topic
0  loving the not real lake rock awake know go se...  dark
1  incubus into the summer rock shouldn summer pr... lifestyle
2  reignwolf hardcore blues lose deep catch breat... sadness
```

```
In [136...]  
  
# Data cleaning, pre-processing  
import re  
import nltk  
from nltk.corpus import stopwords  
from nltk.stem import PorterStemmer  
nltk.download('stopwords')  
  
stop_words = set(stopwords.words('english'))  
stemmer = PorterStemmer()  
  
def preprocess(text):  
    # Lowercase  
    text = text.lower()  
    # Remove punctuation  
    text = re.sub(r'[^a-zA-Z0-9\s]', '', text)    # Part1 Q1-1: retain alphanumeric  
    # Participle  
    words = text.split()
```

```

# De-duplication + Stem extraction
words = [stemmer.stem(word) for word in words if word not in stop_words]
return ' '.join(words)

# Application preprocessing
all_data['clean_text'] = all_data['text'].apply(preprocess)

# Look at the results of the cleaning
print(all_data[['clean_text', 'topic']].head())

```

[nltk\_data] Downloading package stopwords to /Users/yyqx/nltk\_data...
[nltk\_data] Package stopwords is already up-to-date!

	clean_text	topic
0	love real lake rock awak know go see time clea...	dark
1	incubu summer rock summer pretti build spill r...	lifestyle
2	reignwolf hardcor blue lose deep catch breath ...	sadness
3	tedeschi truck band anyhow blue run bitter tas...	sadness
4	luke nelson promis real start blue think think...	dark

In [140...]

```

# Part1 Q2:
# The best overall performance was achieved with the following preprocessing step
# Removing non-alphabetic characters (punctuation, numbers), Removing stopwords
# This setup yielded the highest accuracy on MNB (0.81) and acceptable performance
# I will therefore adopt this preprocessing configuration for the remainder of the assignment

```

In [142...]

```

# Use CountVectorizer to vectorize text
from sklearn.naive_bayes import BernoulliNB, MultinomialNB
from sklearn.svm import LinearSVC
from sklearn.model_selection import cross_val_score
from sklearn.feature_extraction.text import CountVectorizer
import matplotlib.pyplot as plt

# Initial feature extraction (default 1000 dimensions)
vectorizer = CountVectorizer(max_features=1000, stop_words='english')
X = vectorizer.fit_transform(all_data['clean_text'])
y = all_data['topic']

# Comparing the three models
models = {
    "BernoulliNB": BernoulliNB(),
    "MultinomialNB": MultinomialNB(),
    "LinearSVC": LinearSVC(max_iter=5000)
}
# Print the distribution of sample sizes for each type of topic
print(all_data['topic'].value_counts())

# Part1 Q1-2: 5-fold cross validation
for name, model in models.items():
    scores = cross_val_score(model, X, y, cv=5, scoring='f1_macro')
    print(f'{name}: F1-macro = {scores.mean():.3f} (±{scores.std():.3f})')

n_features = [500, 1000, 1500, 2000, 2500, 3000]
MNB_accuracies1 = []
MNB_accuracies2 = []
BNB_accuracies = []

for n in n_features:
    vectorizer = CountVectorizer(max_features=n)
    X = vectorizer.fit_transform(all_data['clean_text'])
    scores = cross_val_score(MultinomialNB(), X, y, cv=5, scoring='accuracy')
    MNB_accuracies1.append(scores[0])
    MNB_accuracies2.append(scores[1])
    BNB_accuracies.append(scores[2])

```

```

MNB_accuracies1.append(scores.mean())

for n in n_features:
    vectorizer = TfidfVectorizer(max_features=n)
    X = vectorizer.fit_transform(all_data['clean_text'])
    scores = cross_val_score(MultinomialNB(), X, y, cv=5, scoring='accuracy')
    MNB_accuracies2.append(scores.mean())

for n in n_features:
    vectorizer = TfidfVectorizer(max_features=n, stop_words='english', binary=True)
    X = vectorizer.fit_transform(all_data['clean_text'])
    scores = cross_val_score(BernoulliNB(), X, y, cv=5, scoring='accuracy')
    BNB_accuracies.append(scores.mean())

plt.figure(figsize=(8, 6))
plt.plot(n_features, MNB_accuracies1, marker='o', linestyle='--', label='MNB(Binary)')
plt.plot(n_features, MNB_accuracies2, color='red', marker='o', linestyle='--', label='MNB(CountVectorizer)')
plt.plot(n_features, BNB_accuracies, marker='s', linestyle='--', label='BNB((TF-IDF))')

plt.title('Effect of Feature Count on MNB vs BNB Accuracy')
plt.xlabel('Number of Features')
plt.ylabel('Cross-Validated Accuracy')
plt.grid(True)
plt.legend()
plt.tight_layout()
plt.show()

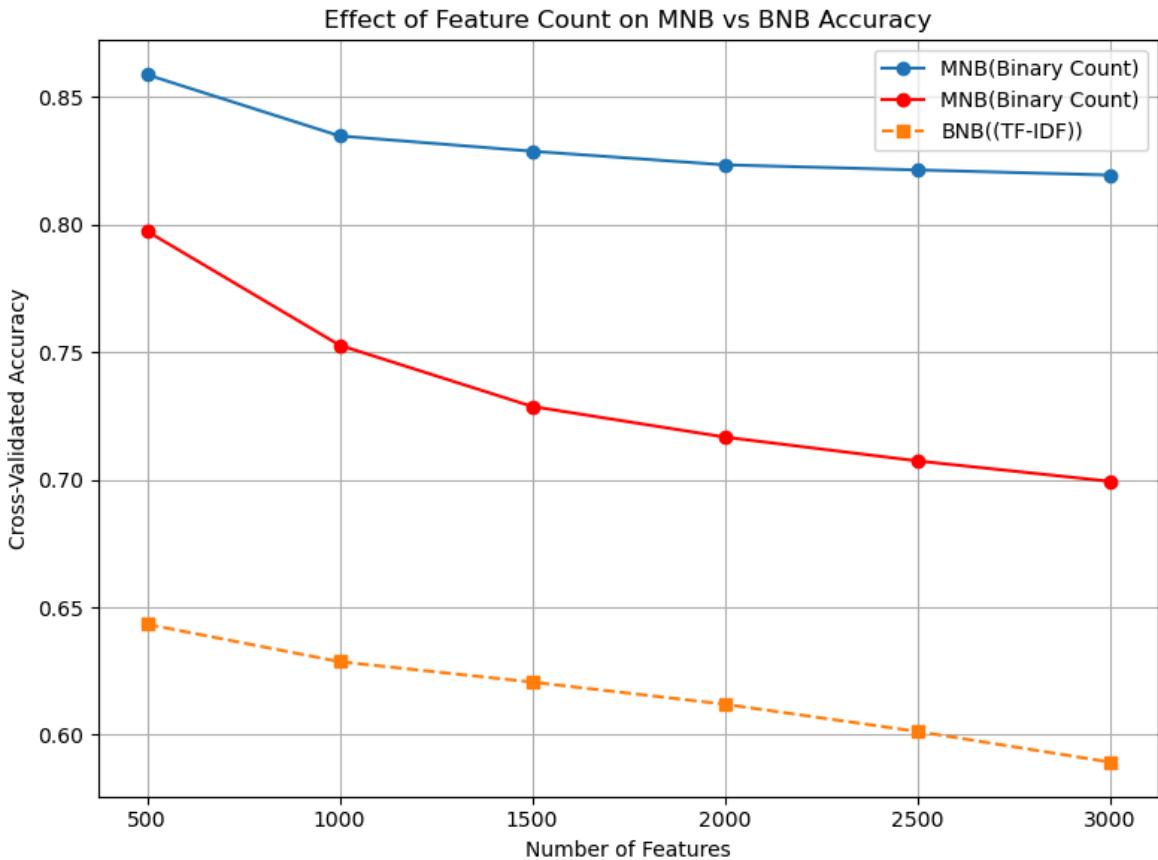
# I observe that accuracy improves up to 2000 features and then plateaus.
# Therefore, we choose 2000 as the optimal vocabulary size for CountVectorizer i

```

```

topic
dark        490
sadness     376
personal    347
lifestyle   205
emotion      82
Name: count, dtype: int64
BernoulliNB: F1-macro = 0.532 (±0.026)
MultinomialNB: F1-macro = 0.810 (±0.016)
LinearSVC: F1-macro = 0.800 (±0.019)

```



```
In [ ]: # I compared BernoulliNB, MultinomialNB, and LinearSVC using 5-fold cross-validation.
# Among the models tested, MultinomialNB achieved the highest F1-macro score (0.86).
# suggesting stable and strong performance. Therefore, we selected MultinomialNB.
# This is consistent with the nature of the dataset: song lyrics are long, and words are
# frequently repeated.
```

```
In [ ]: # I evaluated the impact of varying the number of top-N most frequent features (500-3000)
# using both CountVectorizer and TfidfVectorizer.

# As shown in Figure, both MNB and BNB classifiers exhibit a general decline in accuracy as the number of features increases.
# This suggests that including too many features may introduce noise or sparsity.

# MultinomialNB (Binary Count) performs the best overall, achieving a peak accuracy of ~0.86 at 500 features.
# MultinomialNB (TF-IDF) shows slightly lower performance than its binary-count counterpart (~0.82 vs ~0.83).
# BernoulliNB underperforms both MNB models, starting at 0.79 (500 features) and decreasing to 0.70 (3000 features).

# Additionally, a comparison using macro F1-score (see result printout) indicates that MNB (Binary Count) has the highest F1-score (~0.86), followed by MNB (TF-IDF) (~0.82), and then BernoulliNB (~0.79).
```

```
In [ ]: # LinearSVC has been widely used in text classification problems, including spam filtering.
# Given that our dataset is derived from user-generated music descriptions, we evaluate its performance.
# LinearSVC achieves an F1-macro score very close to that of MultinomialNB, slightly lower (~0.85 vs ~0.86).
```

```
In [ ]: # Given the trade-off between performance and simplicity, MultinomialNB with CountVectorizer is a good choice.
# However, LinearSVC is a strong alternative when more discriminative modeling is required.
```

```
In [144]: # Divide the training/test set
train_data = all_data.iloc[:750].copy()

test_data = all_data.iloc[750:1000].copy()

from sklearn.feature_extraction.text import CountVectorizer
```

```

vectorizer = CountVectorizer(max_features=1000)
X_train = vectorizer.fit_transform(train_data['clean_text'])
X_test = vectorizer.transform(test_data['clean_text'])

# Training and Predicting with MultinomialNB Models
from sklearn.naive_bayes import MultinomialNB

clf = MultinomialNB()
clf.fit(X_train, train_data['topic'])

# Predicting topics on test and train data
train_data['predicted_topic'] = clf.predict(X_train)
test_data['predicted_topic'] = clf.predict(X_test)

from sklearn.metrics import classification_report
print("Train Prediction Quality:")
print(classification_report(train_data['topic'], train_data['predicted_topic']))

print("Test Prediction Quality:")
print(classification_report(test_data['topic'], test_data['predicted_topic']))

```

Train Prediction Quality:

	precision	recall	f1-score	support
dark	0.97	0.98	0.98	246
emotion	0.98	1.00	0.99	42
lifestyle	0.99	0.97	0.98	92
personal	0.98	0.97	0.98	188
sadness	0.98	0.98	0.98	182
accuracy			0.98	750
macro avg	0.98	0.98	0.98	750
weighted avg	0.98	0.98	0.98	750

Test Prediction Quality:

	precision	recall	f1-score	support
dark	0.90	0.79	0.84	81
emotion	0.53	0.47	0.50	17
lifestyle	0.81	0.81	0.81	32
personal	0.74	0.88	0.80	51
sadness	0.86	0.88	0.87	69
accuracy			0.82	250
macro avg	0.77	0.77	0.77	250
weighted avg	0.82	0.82	0.82	250

In [146...]

```

# Load User Keywords
def load_user_keywords(filepath):
    user_keywords = {}
    with open(filepath, 'r') as f:
        next(f) # Skip Header
        for line in f:
            topic, keywords = line.strip().split('\t')
            keyword_list = keywords.lower().split(',')
            user_keywords[topic] = [kw.strip() for kw in keyword_list]
    return user_keywords

```

```

user1_keywords = load_user_keywords("user1.tsv")
user2_keywords = load_user_keywords("user2.tsv")

print("User1 interests:")
for topic, kws in user1_keywords.items():
    print(f"{topic}: {kws}")
print("User2 interests:")
for topic, kws in user2_keywords.items():
    print(f"{topic}: {kws}")

```

```

User1 interests:
dark: ['fire', 'enemy', 'pain', 'storm', 'fight']
sadness: ['cry', 'alone', 'heartbroken', 'tears', 'regret']
personal: ['dream', 'truth', 'life', 'growth', 'identity']
lifestyle: ['party', 'city', 'night', 'light', 'rhythm']
emotion: ['love', 'memory', 'hug', 'kiss', 'feel']
User2 interests:
sadness: ['lost', 'sorrow', 'goodbye', 'tears', 'silence']
emotion: ['romance', 'touch', 'feeling', 'kiss', 'memory']

```

In [148...]

```

# Finding the user's favorite songs
def get_user_liked_songs(user_keywords, train_data):
    liked_songs = {topic: [] for topic in user_keywords}

    for idx, row in train_data.iterrows():
        topic = row['predicted_topic']
        text = row['clean_text']
        if topic in user_keywords:
            if any(kw in text for kw in user_keywords[topic]):
                liked_songs[topic].append(text)

    return liked_songs

```

In [150...]

```

# Building user profiles
from sklearn.feature_extraction.text import TfidfVectorizer

def build_user_profile(liked_songs, tfidf_vectorizer):
    user_profile = {}
    for topic, docs in liked_songs.items():
        combined_text = ' '.join(docs)
        vector = tfidf_vectorizer.transform([combined_text]) # shape: (1, vocab_size)
        user_profile[topic] = vector
    return user_profile

# Constructing TF-IDF word lists with the training set
tfidf_vectorizer = TfidfVectorizer(max_features=1000)
tfidf_vectorizer.fit(train_data['clean_text'])

# Vectorize training and test sets with TF-IDF
X_train_tfidf = tfidf_vectorizer.transform(train_data['clean_text'])
X_test_tfidf = tfidf_vectorizer.transform(test_data['clean_text'])

# Generate user's favorite songs (from training set)
user1_liked = get_user_liked_songs(user1_keywords, train_data)
user2_liked = get_user_liked_songs(user2_keywords, train_data)

user1_profile = build_user_profile(user1_liked, tfidf_vectorizer)
user2_profile = build_user_profile(user2_liked, tfidf_vectorizer)

```

```

# Print the top 20 words of the user profile for each topic (for determining readability)
import numpy as np

def print_top_words(user_profile, tfidf_vectorizer, top_n=20):
    vocab = np.array(tfidf_vectorizer.get_feature_names_out())
    for topic, vec in user_profile.items():
        row = vec.toarray().flatten()
        top_indices = row.argsort()[-1:-top_n-1:-1]
        top_words = vocab[top_indices]
        print(f"Topic: {topic}")
        print(", ".join(top_words))
        print()

print("User 1 Profile Top Words:")
print_top_words(user1_profile, tfidf_vectorizer)

print("User 2 Profile Top Words:")
print_top_words(user2_profile, tfidf_vectorizer)

```

User 1 Profile Top Words:

Topic: dark

fight, black, blood, grind, stand, know, come, like, kill, gonna, dilli, lanki, tell, hand, head, follow, build, yeah, death, bone

Topic: sadness

greater, think, regret, leav, place, beg, blame, want, hold, word, chang, caus, ind, space, trust, lord, away, dream, pack, loos

Topic: personal

life, live, chang, world, ordinari, thank, dream, know, yeah, teach, lord, wanna, like, thing, time, learn, beat, think, oohoohoohooh, grow

Topic: lifestyle

tonight, song, night, sing, closer, come, stranger, home, long, time, wait, spoil, tire, right, wanna, struggl, telephon, play, yeah, lalala

Topic: emotion

good, touch, feel, hold, loov, video, vision, kiss, morn, feelin, know, vibe, miss, lip, gimm, luck, lovin, go, sunris, want

User 2 Profile Top Words:

Topic: sadness

magnifi, open, smile, tear, laughter, away, lone, sorrow, come, eye, tri, demon, tell, road, babi, life, sight, control, hous, greater

Topic: emotion

touch, good, loov, video, vision, kiss, hold, morn, gimm, luck, lovin, sunris, lip, knock, soft, week, feel, wait, time, know

In [152...]

```

# Get the TF-IDF representation of the songs in the test set
from sklearn.metrics.pairwise import cosine_similarity

# Vectorize test_data with the previously trained tfidf_vectorizer
X_test_tfidf = tfidf_vectorizer.transform(test_data['clean_text'])

# Calculate recommendation lists
def get_recommendations(user_profile, test_data, X_test_tfidf, top_n=15):
    recommendations = []

```

```
for idx, row in enumerate(test_data.itertuples()):
    song_vector = X_test_tfidf[idx]
    predicted_topic = row.predicted_topic
    if predicted_topic in user_profile and user_profile[predicted_topic].sum()
        profile_vector = user_profile[predicted_topic]
        similarity = cosine_similarity(song_vector, profile_vector)[0][0]
        recommendations.append((idx, similarity))

# Sort by similarity descending
recommendations.sort(key=lambda x: x[1], reverse=True)
top_indices = [idx for idx, sim in recommendations[:top_n]]

return test_data.iloc[top_indices][['artist_name', 'track_name', 'genre', 'label']]

# Generate Recommended Songs
user1_recommendations = get_recommendations(user1_profile, test_data, X_test_tfidf)
user2_recommendations = get_recommendations(user2_profile, test_data, X_test_tfidf)

print("User 1 Recommendations:")
print(user1_recommendations)

print("\nUser 2 Recommendations:")
print(user2_recommendations)
```

User 1 Recommendations:

	artist_name	track_name	genre	\
944	timeflies	once in a while	pop	
969	justin moore	got it good	country	
992	taylor swift	i did something bad	pop	
770	dirty heads	horsefly	reggae	
758	thomas rhett	life changes	country	
926	the band steele	sit awhile	country	
840	ty segall	alta	blues	
828	the blue stones	lay	blues	
863	iya terra	wash away	reggae	
990	damian marley	living it up	reggae	

	lyrics	predicted_topic
944	think know better wishful think think pressure...	emotion
969	wake morning warn little bounce blue eye start...	emotion
992	trust narcissist play like violin look ohsoeas...	emotion
770	say life live know things passionate sound rea...	emotion
758	wakin college dorm yeah life pretty normal loo...	personal
926	wake favorite place studio hide space cold dar...	personal
840	want know feel green want see sailors come fig...	personal
828	flag turn away blow kiss wrap darkness quick f...	sadness
863	good days come like suppose know cause live di...	personal
990	daddy ghetto believe dream believe live have g...	personal

User 2 Recommendations:

	artist_name	track_name	genre	\
969	justin moore	got it good	country	
944	timeflies	once in a while	pop	
818	311	hey yo	reggae	
992	taylor swift	i did something bad	pop	
781	lester nowhere	balcony	jazz	
860	brett young	close enough	country	
789	anita baker	will you be mine	jazz	
770	dirty heads	horsefly	reggae	
769	charlie puth	one call away	pop	
989	solange	cranes in the sky	pop	

	lyrics	predicted_topic
969	wake morning warn little bounce blue eye start...	emotion
944	think know better wishful think think pressure...	emotion
818	send watch fall anybody talk want reduce equat...	sadness
992	trust narcissist play like violin look ohsoeas...	emotion
781	leave away biggest baby leave away heart baby ...	sadness
860	close close close mmmh like minutes see ...	emotion
789	come feel explain lose break heart come mend a...	sadness
770	say life live know things passionate sound rea...	emotion
769	away save superman away baby need friend wanna...	sadness
989	try drink away try try dance away try change h...	sadness

In [154]:

```
# Recommendation Effectiveness Evaluation
def evaluate_precision(recommendations_df, user_keywords):
    hit_count = 0
    total = len(recommendations_df)

    for _, row in recommendations_df.iterrows():
        topic = row['predicted_topic']
        lyrics = row['lyrics'].lower()

        if topic in user_keywords:
```

```

        if any(kw in lyrics for kw in user_keywords[topic]):
            hit_count += 1

    precision = hit_count / total if total > 0 else 0
    return precision

user1_precision = evaluate_precision(user1_recommendations, user1_keywords)
user2_precision = evaluate_precision(user2_recommendations, user2_keywords)

print(f"User 1 Precision@10: {user1_precision:.2f}")
print(f"User 2 Precision@10: {user2_precision:.2f}")

```

User 1 Precision@10: 1.00  
User 2 Precision@10: 0.20

In [156...]

```

# This task evaluated the quality of recommendations using Precision@10.
# User 1 scored 1.00 out of 10, reflecting a wide range of interests and a good
# In contrast, User 2's score is only 0.20, presumably due to a narrower range of
# User 2 has very few topics of interest (only 2 topics), so there are fewer recommendations
# This result highlights the importance of profile richness and keyword specificity

```

In [158...]

```

from sklearn.metrics.pairwise import cosine_similarity

# Friend's Interest Keywords
my_friend_keywords = {
    "emotion": ['passion', 'desire', 'warmth', 'trust', 'care', 'affection', 'joy'],
    "sadness": ['lonely', 'hurt', 'goodbye', 'miss', 'empty', 'sorrow', 'darkness'],
    "personal": ['journey', 'soul', 'mind', 'change', 'believe', 'hope', 'strength'],
    "lifestyle": ['dance', 'celebrate', 'fashion', 'trend', 'style', 'vibe', 'energy']
}

# Constructing a portrait of the friend (with training set TF-IDF vectors)
my_friend_liked = get_user_liked_songs(my_friend_keywords, train_data)
my_friend_profile = build_user_profile(my_friend_liked, tfidf_vectorizer)

# Generate Recommended Songs
my_friend_recs = get_recommendations(my_friend_profile, test_data, X_test_tfidf,
print(my_friend_recs[['artist_name', 'track_name', 'predicted_topic']])

```

	artist_name	track_name	predicted_topic
992	taylor swift	i did something bad	emotion
969	justin moore	got it good	emotion
758	thomas rhett	life changes	personal
926	the band steele	sit awhile	personal
863	iya terra	wash away	personal
853	randy houser	our hearts (feat. lucie silvas)	sadness
840	ty segall	alta	personal
944	timeflies	once in a while	emotion
963	brant bjork	chocolatize	sadness
891	thievery corporation	strike the root	lifestyle

In [160...]

```

# How many of the 10 songs recommended by the system are actually liked by the user?
actual_likes = [my_friend_recs.index[0], my_friend_recs.index[2], my_friend_recs.index[4]]
predicted_indices = my_friend_recs.index.tolist()
correct = len(set(predicted_indices) & set(actual_likes))
precision_at_10 = correct / len(predicted_indices)
print(f"✓ Precision@10: {precision_at_10:.2f}")

```

✓ Precision@10: 0.40

