

## zid:z5598515 name: Hang Pan

```
In [341... import pandas as pd
```

```
#Load data
df=pd.read_csv("./data/dataset.tsv",sep="\t", encoding="utf-8")
```

```
df
```

```
Out[343... 
```

	artist_name	track_name	release_date	genre	lyrics	topic
0	loving	the not real lake	2016	rock	awake know go see time clear world mirror worl...	dark
1	incubus	into the summer	2019	rock	shouldn summer pretty build spill ready overfl...	lifestyle
2	reignwolf	hardcore	2016	blues	lose deep catch breath think say try break wal...	sadness
3	tedeschi trucks band	anyhow	2016	blues	run bitter taste take rest feel anchor soul pl...	sadness
4	lukas nelson and promise of the real	if i started over	2017	blues	think think different set apart sober mind sym...	dark
...	...	...	...	...	...	...
1495	ra ra riot	absolutely	2016	rock	year absolutely absolutely absolutely crush ab...	emotion
1496	mat kearney	face to face	2018	rock	breakthrough hours hear truth moments trade fa...	dark
1497	owane	born in space	2018	jazz	look look right catch blue eye own state breat...	dark
1498	nappy roots	blowin' trees	2019	hip hop	nappy root gotta alright flyin dear leave lone...	personal
1499	skillet	stars	2016	rock	speak word life begin tell oceans start motion...	sadness

1500 rows × 6 columns

```
In [344... # combine_field
def combine_fields(row):
    return f'{row['artist_name']} {row['track_name']} {row['genre']} {row['lyric']}
```

```
In [345... new_df=pd.DataFrame({'Content': df['text'], 'Category': df['topic']})
```

```
In [346... #do data cleaning
# Drop duplicates and missing values
new_df = new_df.drop_duplicates()
new_df = new_df.dropna()
```

```
In [347... new_df.head()
```

Out[347... 

	Content	Category
<b>0</b>	loving the not real lake rock awake know go se...	dark
<b>1</b>	incubus into the summer rock shouldn summer pr...	lifestyle
<b>2</b>	reignwolf hardcore blues lose deep catch breat...	sadness
<b>3</b>	tedeschi trucks band anyhow blues run bitter t...	sadness
<b>4</b>	lukas nelson and promise of the real if i star...	dark

## Part1

### Q1.

```
In [348... import nltk
import re
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import PorterStemmer

nltk.download('stopwords')
nltk.download('punkt_tab')
nltk.download('punkt')
ps = PorterStemmer()
stop_words = set(stopwords.words('english'))

# Define preprocessing function
def preprocess_text(text):
    text = text.lower()
    # text = re.sub(r'^\w\s]', '', text)      # Check what this removes --- might
    text = re.sub(r'\s+', ' ', text)
    text=re.sub(r"^\w\s'\-]", ' ', text)
    text = re.sub(r'\s+', ' ', text).strip()           # Remove extra space
    tokens = word_tokenize(text)
    tokens = [word for word in tokens if word not in stop_words]
    tokens = [ps.stem(word) for word in tokens]
    return ' '.join(tokens)
```

```
# Apply preprocessing to each document
new_df['Content'] = new_df['Content'].apply(preprocess_text)
```

```
[nltk_data] Downloading package stopwords to
[nltk_data]      /Users/russell/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt_tab to
[nltk_data]      /Users/russell/nltk_data...
[nltk_data] Package punkt_tab is already up-to-date!
[nltk_data] Downloading package punkt to /Users/russell/nltk_data...
[nltk_data] Package punkt is already up-to-date!
```

```
In [349]: print(new_df['Content'].head())
```

```
0    love real lake rock awak know go see time clea...
1    incubu summer rock summer pretti build spill r...
2    reignwolf hardcor blue lose deep catch breath ...
3    tedeschi truck band anyhow blue run bitter tas...
4    luka nelson promis real start blue think think...
Name: Content, dtype: object
```

```
In [350]: from sklearn.model_selection import cross_val_score
```

```
from sklearn.pipeline import Pipeline
from sklearn.naive_bayes import BernoulliNB
from sklearn.feature_extraction.text import CountVectorizer
# from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.model_selection import StratifiedKFold
```

```
skf = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)
```

```
# Create pipeline: vectorizer + classifier
```

```
bnb_pipeline = Pipeline([
    ('vectorizer', CountVectorizer()),
    ('classifier', BernoulliNB())
])
```

```
# Evaluate with 5-fold cross-validation
```

```
scores = cross_val_score(bnb_pipeline, new_df['Content'], new_df['Category'], cv=5)
print("BernoulliNB Cross-Validated Macro F1 Score: {:.4f}".format(scores.mean()))
```

```
BernoulliNB Cross-Validated Macro F1 Score: 0.3398
```

```
In [351]: from sklearn.model_selection import cross_val_predict
from sklearn.metrics import classification_report
```

```
y_pred_cv = cross_val_predict(bnb_pipeline, new_df['Content'], new_df['Category'])
print(classification_report(new_df['Category'], y_pred_cv))
```

	precision	recall	f1-score	support
dark	0.64	0.77	0.70	487
emotion	0.00	0.00	0.00	79
lifestyle	0.00	0.00	0.00	202
personal	0.68	0.32	0.43	341
sadness	0.43	0.84	0.56	371
accuracy			0.54	1480
macro avg	0.35	0.39	0.34	1480
weighted avg	0.47	0.54	0.47	1480

```
/Users/russell/miniconda3/envs/python39/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
/Users/russell/miniconda3/envs/python39/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
/Users/russell/miniconda3/envs/python39/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

## Part 1 - Q1: Addressing Simplifications in the Tutorial

The tutorial notebook contained two key simplifications:

### 1. Overly aggressive regex in preprocessin

The original regular expression `re.sub(r'^\w\s]', ' ', text)` removes all non-word characters, including potentially meaningful ones such as apostrophes (' '), hyphens ,It is more suitable for English corpora in natural language processing and avoids destroying the structure of words(-), or musical symbols that may help understand genre-specific lyrics and song strucdre.

**Fix:** We modified the regex to the following:

```
re.sub(r"^\w\s'\-]", ' ', text)
```

### 2. Using cross-validation

The original tutorial uses a single random train-test split, which can result in unstable and unrepresentative evaluation due to sampling variance. StratifiedKFold ensures that in each compromise, the Category distribution of new\_df['Category'] is consistent with the entire dataset **Fix:** We replaced this with 5-fold cross-validation using `cross_val_score` from `sklearn.model_selection`:

```
scores = cross_val_score(pipeline, X, y, cv=sfk, scoring='accuracy')
```

We used `cross_val_score` with the default `scoring='accuracy'`, but we note that:

- Accuracy is suitable only when classes are balanced.
- In later questions, we will include other metrics such as macro-F1 and weighted-F1 which are better for class imbalance (as seen in the lecture: macro-averaging is less dominated by large classes). we just do experiment here.

## Q2.

### Multinomial Naive Bayes

In [352...]

```
import re
import nltk
from nltk.corpus import stopwords
from nltk.tokenize import word_tokenize
from nltk.stem import PorterStemmer, WordNetLemmatizer
from sklearn.pipeline import Pipeline
from sklearn.model_selection import cross_val_score
from sklearn.naive_bayes import MultinomialNB
from sklearn.feature_extraction.text import CountVectorizer

nltk.download('stopwords')
nltk.download('punkt')
nltk.download('wordnet')

stop_words = set(stopwords.words('english'))
ps = PorterStemmer()
lemmatizer = WordNetLemmatizer()

# ----- Different Preprocessing Versions -----

def preprocess_v1(text):
    # simple Lowercasing and tokenization, no stopword removal
    text = text.lower()
    text = re.sub(r"[\w\s'-]", ' ', text)
    text = re.sub(r"\s+", " ", text).strip()
    tokens = word_tokenize(text)
    return ' '.join(tokens)

def preprocess_v2(text):
    # simple Lowercasing and tokenization + stopword removal
    text = text.lower()
    text = re.sub(r"[\w\s'-]", ' ', text)
    text = re.sub(r"\s+", " ", text).strip()
    tokens = word_tokenize(text)
    tokens = [word for word in tokens if word not in stop_words]
    return ' '.join(tokens)

def preprocess_v3(text):
    # simple Lowercasing and tokenization + stopword + stemming
    text = text.lower()
    text = re.sub(r"[\w\s'-]", ' ', text)
    text = re.sub(r"\s+", " ", text).strip()
    tokens = word_tokenize(text)
    tokens = [word for word in tokens if word not in stop_words]
    tokens = [ps.stem(word) for word in tokens]
    return ' '.join(tokens)

def preprocess_v4(text):
    # simple Lowercasing and tokenization + stopword + lemmatization
    text = text.lower()
    text = re.sub(r"[\w\s'-]", ' ', text)
    text = re.sub(r"\s+", " ", text).strip()
    tokens = word_tokenize(text)
    tokens = [word for word in tokens if word not in stop_words]
    tokens = [lemmatizer.lemmatize(word) for word in tokens]
    return ' '.join(tokens)

from sklearn.metrics import classification_report
```

```

def evaluate_preprocessing(preprocess_func, name, use_sklearn_stopwords=False):
    processed_text = new_df['Content'].apply(preprocess_func)
    sfk = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

    if use_sklearn_stopwords:
        vectorizer = CountVectorizer(stop_words='english')
    else:
        vectorizer = CountVectorizer()

    pipeline = Pipeline([
        ('vectorizer', vectorizer),
        ('classifier', MultinomialNB())
    ])

    f1_scores = []
    acc_scores = []
    for train_idx, test_idx in sfk.split(processed_text, new_df['Category']):
        X_train, X_test = processed_text.iloc[train_idx], processed_text.iloc[test_idx]
        y_train, y_test = new_df['Category'].iloc[train_idx], new_df['Category'].iloc[test_idx]
        pipeline.fit(X_train, y_train)
        y_pred = pipeline.predict(X_test)
        report = classification_report(y_test, y_pred, output_dict=True, zero_division=1)
        acc_scores.append(report['accuracy'])
        f1_scores.append(report['macro avg']['f1-score'])

    print(f'{name}: accuracy = {sum(acc_scores)/5:.4f}, macro F1 = {sum(f1_scores)/5:.4f}')
    return {'accuracy': sum(acc_scores)/5, 'macro_f1': sum(f1_scores)/5}

```

```

results = {}
results['V1'] = evaluate_preprocessing(preprocess_v1, 'V1 - Lowercase only')
results['V2'] = evaluate_preprocessing(preprocess_v2, 'V2 - +Stopwords removed')
results['V3'] = evaluate_preprocessing(preprocess_v3, 'V3 - +Stemming')
results['V4'] = evaluate_preprocessing(preprocess_v4, 'V4 - +Lemmatization')
results['V5'] = evaluate_preprocessing(preprocess_v1, 'V5 - sklearn stopwords', True)

# ----- Summary -----
import pandas as pd
import matplotlib.pyplot as plt

summary_df = pd.DataFrame.from_dict(results, orient='index', columns=['accuracy'])
summary_df = summary_df.sort_values(by='accuracy', ascending=False)
print(summary_df)

summary_df.plot(kind='bar', legend=False, title='Preprocessing Strategy Comparison')
plt.xticks(rotation=45)
plt.tight_layout()
plt.show()

```

```


[nltk_data] Downloading package stopwords to
[nltk_data]      /Users/russell/nltk_data...
[nltk_data] Package stopwords is already up-to-date!
[nltk_data] Downloading package punkt to /Users/russell/nltk_data...
[nltk_data] Package punkt is already up-to-date!
[nltk_data] Downloading package wordnet to /Users/russell/nltk_data...
[nltk_data] Package wordnet is already up-to-date!

```

```

V1 - Lowercase only: accuracy = 0.8088, macro F1 = 0.7457
V2 - +Stopwords removed: accuracy = 0.8081, macro F1 = 0.7453
V3 - +Stemming: accuracy = 0.8081, macro F1 = 0.7452
V4 - +Lemmatization: accuracy = 0.8095, macro F1 = 0.7463
V5 - sklearn stopwords: accuracy = 0.8027, macro F1 = 0.7331
    accuracy
V4 0.809459
V1 0.808784
V2 0.808108
V3 0.808108
V5 0.802703

```



## Summary

To identify the most effective preprocessing strategy for topic classification using Multinomial Naive Bayes with CountVectorizer, I evaluated five variations:

Preprocessing Version	Description	Accuracy
V4	Lowercase + NLTK stopwords + Lemmatization	0.8088
V1	Lowercase only	0.8081
V2	Lowercase + NLTK stopwords	0.8081
V3	Lowercase + NLTK stopwords + Stemming	0.8081
V5	Lowercase + sklearn stopwords	0.8034

## Conclusion:

Version 4 (lemmatization) achieved the highest accuracy and F1-score. This suggests that reducing words to their base forms using `WordNetLemmatizer` is more beneficial than

crude stemming, and outperforms simpler strategies like just lowercasing or removing stopwords.

For the rest of the assignment, I will use the V4 preprocessing strategy as my standard pipeline.

### Q3.


## Check whether the distribution of sample categories is uniform

In [353...]

```
import matplotlib.pyplot as plt

category_counts = new_df['Category'].value_counts()

plt.figure(figsize=(8, 5))
category_counts.plot(kind='bar', color='skyblue', edgecolor='black')
plt.title('Category Distribution')
plt.xlabel('Category')
plt.ylabel('Count')
plt.xticks(rotation=45)
plt.grid(axis='y', linestyle='--', alpha=0.7)
plt.tight_layout()
plt.show()
```



In [354...]

```
from sklearn.metrics import classification_report
from sklearn.model_selection import cross_val_predict
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import BernoulliNB, MultinomialNB

# The preprocessed text by v4
X = new_df['Content'].apply(preprocess_v4)
```

```

y = new_df['Category']
sfk = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# BernoulliNB
bnb_pipeline = Pipeline([
    ('vectorizer', CountVectorizer()),
    ('classifier', BernoulliNB())
])
bnb_preds = cross_val_predict(bnb_pipeline, X, y, cv=sfk)
bnb_report = classification_report(y, bnb_preds, output_dict=True)

# MultinomialNB
mnb_pipeline = Pipeline([
    ('vectorizer', CountVectorizer()),
    ('classifier', MultinomialNB())
])
mnb_preds = cross_val_predict(mnb_pipeline, X, y, cv=sfk)
mnb_report = classification_report(y, mnb_preds, output_dict=True)

```

```

/Users/russell/miniconda3/envs/python39/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and
being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/Users/russell/miniconda3/envs/python39/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and
being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
/Users/russell/miniconda3/envs/python39/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and
being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))

```

In [355...]

```

import pandas as pd


df_compare = pd.DataFrame({
    'BNB': {
        'accuracy': bnb_report['accuracy'],
        'macro f1': bnb_report['macro avg']['f1-score'],
        'macro precision': bnb_report['macro avg']['precision'],
        'weighted f1': bnb_report['weighted avg']['f1-score'],
    },
    'MNB': {
        'accuracy': mnb_report['accuracy'],
        'macro f1': mnb_report['macro avg']['f1-score'],
        'macro precision': mnb_report['macro avg']['precision'],
        'weighted f1': mnb_report['weighted avg']['f1-score'],
    }
})

print(df_compare)
df_compare.T.plot(kind='bar', title='BNB vs MNB Performance Comparison', ylabel=

```

	BNB	MNB
accuracy	0.539189	0.809459
macro f1	0.340520	0.747706
macro precision	0.349283	0.827808
weighted f1	0.473170	0.803276

```
Out[355]: <Axes: title={'center': 'BNB vs MNB Performance Comparison'}, ylabel='Score'>
```



We evaluated two models - Bernoulli Naive Bayes (BNB) and Multinomial Naive Bayes (MNB), using 5-fold stratified cross-validation on the full dataset. The goal was to evaluate overall classification performance using the metrics discussed in the lecture:

Accuracy: measures the proportion of correct predictions; however, it tends to be biased towards the majority class.

Macro-F1-score: unweighted average of all class F1 scores; treats each class equally, better suited for imbalanced datasets.

Macro-precision: average precision across all classes; helps assess the model's tendency to produce false positives.

Weighted-F1-score: similar to macro-F1, but accounts for class imbalance by weighting the score by class frequency.

Metric	BNB	MNB
Accuracy	0.5392	0.8095
Macro F1-score	0.3405	0.7477
Macro Precision	0.3493	0.8278
Weighted F1-score	0.4732	0.8033

While we report multiple evaluation metrics, we choose the macro F1-score as the main metric. This is because:

- The dataset shows some imbalance between the 5 topic classes.

- Macro F1 gives equal weight to each class, regardless of its frequency, making it more appropriate for assessing performance across majority classes.
- In contrast, accuracy can be dominated by larger classes, and may give a misleading impression of model quality.

Therefore, although we present accuracy,macro-f1,macro-precision and weighted F1 for completeness, we choice the macro F1-score as the main metrics.

Conclusion: MultinomialNB has better outperforms BernoulliNB across all metrics above, especially in macro-averaged scores, indicating significantly better performance across all classes — including minority ones. Therefore, MultinomialNB is chosen as the preferred model for the remainder of this assignment. .

In [356]..

```
from sklearn.model_selection import cross_val_predict
from sklearn.metrics import classification_report
from sklearn.pipeline import Pipeline
from sklearn.feature_extraction.text import CountVectorizer
from sklearn.naive_bayes import BernoulliNB, MultinomialNB
import pandas as pd
import matplotlib.pyplot as plt


X = new_df['Content'].apply(preprocess_v4)
y = new_df['Category']
# find the best params
feature_list = [50,100,300,350,400,500,600,700, 1000, 2000, 5000, 10000]
results = []
sfk = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

for n in feature_list:
    for model_cls in [BernoulliNB, MultinomialNB]:
        model_name = model_cls.__name__
        vectorizer = CountVectorizer(max_features=n)
        pipeline = Pipeline([
            ('vectorizer', vectorizer),
            ('classifier', model_cls())
        ])
        y_pred = cross_val_predict(pipeline, X, y, cv=sfk)
        report = classification_report(y, y_pred, output_dict=True)
        results.append({
            'model': model_name,
            'max_features': n,
            'accuracy': report['accuracy'],
            'macro_f1': report['macro avg']['f1-score'],
            'weighted_f1': report['weighted avg']['f1-score'],
            'macro_precision': report['macro avg']['precision']
        })
df_results = pd.DataFrame(results)
```


```
/Users/russell/miniconda3/envs/python39/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
/Users/russell/miniconda3/envs/python39/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
/Users/russell/miniconda3/envs/python39/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
/Users/russell/miniconda3/envs/python39/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
/Users/russell/miniconda3/envs/python39/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))  
/Users/russell/miniconda3/envs/python39/lib/python3.9/site-packages/sklearn/metrics/_classification.py:1565: UndefinedMetricWarning: Precision is ill-defined and being set to 0.0 in labels with no predicted samples. Use `zero_division` parameter to control this behavior.  
    _warn_prf(average, modifier, f"{metric.capitalize()} is", len(result))
```

```
In [357...]:  
for metric in ['accuracy', 'macro_f1', 'weighted_f1', 'macro_precision']:  
    df_pivot = df_results.pivot(index='max_features', columns='model', values=metric)  
    df_pivot.plot(kind='bar', title=f'{metric.upper()} vs max_features', ylabel=metric)  
    plt.xticks(rotation=0)  
    plt.tight_layout()  
    plt.show()
```


ACCURACY vs max\_features




MACRO\_F1 vs max\_features



WEIGHTED\_F1 vs max\_features



MACRO\_PRECISION vs max\_features



In [362]:

df\_results

Out[362...]

	model	max_features	accuracy	macro_f1	weighted_f1
<b>0</b>	BernoulliNB	100	0.572297	0.490605	0.565003
<b>1</b>	MultinomialNB	100	0.754730	0.730993	0.754451
<b>2</b>	BernoulliNB	300	0.635811	0.549783	0.630718
<b>3</b>	MultinomialNB	300	0.862162	0.846457	0.862045
<b>4</b>	BernoulliNB	400	0.642568	0.550424	0.636851
<b>5</b>	MultinomialNB	400	0.862838	0.842607	0.862987
<b>6</b>	BernoulliNB	500	0.643919	0.544675	0.637370
<b>7</b>	MultinomialNB	500	0.874324	0.853545	0.874410
<b>8</b>	BernoulliNB	600	0.642568	0.541802	0.634626
<b>9</b>	MultinomialNB	600	0.868919	0.844918	0.868894
<b>10</b>	BernoulliNB	1000	0.630405	0.517844	0.617114
<b>11</b>	MultinomialNB	1000	0.849324	0.833164	0.849234
<b>12</b>	BernoulliNB	2000	0.620946	0.487596	0.599266
<b>13</b>	MultinomialNB	2000	0.837838	0.813920	0.837261
<b>14</b>	BernoulliNB	5000	0.567568	0.391704	0.519969
<b>15</b>	MultinomialNB	5000	0.820946	0.775257	0.817672
<b>16</b>	BernoulliNB	10000	0.539189	0.340520	0.473170
<b>17</b>	MultinomialNB	10000	0.809459	0.747706	0.803276

In order to, studying the effect of vocabulary size on classification performance, we varied the max\_features parameter in CountVectorizer, limiting the number of most frequent words used for feature extraction. We tested values from 100 to 10,000 across both BernoulliNB and MultinomialNB classifiers, using 5-fold stratified cross-validation.

we choice three metrics:

- Accuracy: Overall correctness.
- Macro F1-score: Preferred metric (see Q3), treats all classes equally.
- Weighted F1-score: Adjusts for class imbalance, favors majority class performance.

For MultinomialNB, performance peaked around at 500, where:

- Accuracy = 0.8743
- Macro F1 = 0.8535
- Weighted F1 = 0.8744
- Beyond 1000, macro-F1 declined, indicating possible overfitting or noise from infrequent tokens.

So, we choose max\_features = 500 and MultinomialNB for the remainder of the assignment. This configuration gave the best overall performance across all metrics, while maintaining model simplicity.

## Q5.

Here, we choose the LogisticRegression as new model over MNB. This is because, Logistic Regression is a linear classifier algorithm that describes probability of class membership by way of logistic (sigmoid) function. Logistic Regression works well for high-dimensional, sparse feature spaces like in text classification task with bag-of-words. Since song lyrics contain many words but few per text, Logistic Regression is particularly suited to handling this type of sparsity. It also handles correlation between features (i.e., the co-occurrence of words) well, and so is a good option for topic classification in this data.

In [366...]

```
from sklearn.linear_model import LogisticRegression
from sklearn.model_selection import GridSearchCV, cross_val_predict
from sklearn.pipeline import Pipeline
from sklearn.metrics import classification_report
from sklearn.feature_extraction.text import CountVectorizer

X = new_df['Content'].apply(preprocess_v4)
y = new_df['Category']

# Pipeline with Logistic Regression and Grid Search
logreg_pipeline = Pipeline([
    ('vectorizer', CountVectorizer(max_features=500)),
    ('classifier', LogisticRegression(max_iter=1000))
])
sfk = StratifiedKFold(n_splits=5, shuffle=True, random_state=42)

# Hyperparameter grid search
param_grid = {
    'classifier__C': [0.01, 0.1, 1, 10]
}

grid = GridSearchCV(logreg_pipeline, param_grid, cv=sfk, scoring='f1_macro', n_jobs=-1)
grid.fit(X, y)

# Optimal parameters
print("Best C:", grid.best_params_)
# Save best model
best_model = grid.best_estimator_

# Make predictions using the best model
y_pred = cross_val_predict(best_model, X, y, cv=sfk)
report = classification_report(y, y_pred, output_dict=True)
```

Best C: {'classifier\_\_C': 0.1}

In [367...]

```
comparison = pd.DataFrame({
    'BNB': {
        'accuracy': bnb_report['accuracy'],
        'macro_f1': bnb_report['macro avg']['f1-score'],
        'weighted_f1': bnb_report['weighted avg']['f1-score']
    },
    'MNB': {
```

```


        'accuracy': mnb_report['accuracy'],
        'macro_f1': mnb_report['macro avg']['f1-score'],
        'weighted_f1': mnb_report['weighted avg']['f1-score']
    },
    'LogReg': {
        'accuracy': report['accuracy'],
        'macro_f1': report['macro avg']['f1-score'],
        'weighted_f1': report['weighted avg']['f1-score']
    }
})

print(comparison)
comparison.T.plot(kind='bar', title='BNB vs MNB vs Logistic Regression', ylabel=

```

	BNB	MNB	LogReg
accuracy	0.539189	0.809459	0.868243
macro_f1	0.340520	0.747706	0.854881
weighted_f1	0.473170	0.803276	0.867571

Out[367... <Axes: title={'center': 'BNB vs MNB vs Logistic Regression'}, ylabel='Score'>



## Q5: Comparing Logistic Regression with Naive Bayes

We used `sklearn.linear_model.LogisticRegression` with the following settings:

- `max_iter = 1000` to ensure convergence
- `CountVectorizer(max_features=500)` as determined in Q4
- Preprocessing: lowercase, stopword removal, and lemmatization (from Q2 V4)

Metric	BNB	MNB	Logistic Regression
Accuracy	0.5392	0.8095	0.8669
Macro F1	0.3405	0.7477	0.8527
Weighted F1	0.4732	0.8033	0.8662

Logistic Regression outperforms both BNB and MNB across all metrics, especially in macro and weighted F1, indicating better performance across all classes, including minority ones.

#### Hypothesis and Result:

We hypothesized that Logistic Regression would outperform both BernoulliNB and MultinomialNB, given its ability to model feature weights directly and its robust performance in text classification tasks.

This hypothesis was confirmed by experimental results. Logistic Regression achieved the highest scores across all key evaluation metrics:

Accuracy: 0.8669

Macro F1: 0.8527

Weighted F1: 0.8662

In contrast, MultinomialNB and BernoulliNB performed significantly worse, especially in macro-averaged scores, which are crucial for imbalanced datasets. These results indicate that Logistic Regression is the most effective method for topic classification on this dataset.

## Part2

```
In [368...]: import pandas as pd
In [369...]: from sklearn.model_selection import train_test_split
In [370...]: df_part2=pd.read_csv("./data/dataset.tsv",sep="\t", encoding="utf-8")
In [371...]: df_part2 = df_part2.drop_duplicates()
           df_part2 = df_part2.dropna()
In [372...]: def combine_song_fields(row):
              return f'{row['artist_name']} {row['track_name']} {row['genre']} {row['lyric']}
df_part2['full_text'] = df_part2.apply(combine_song_fields, axis=1)
```

#### use best model to predict (LR)

```
In [373...]: df_part2['processed'] = df_part2['full_text'].apply(preprocess_v4)
#use LR model to predict which save at before
df_part2['predicted_topic'] = best_model.predict(df_part2['processed'])
```

```
In [374...]: train_df = df_part2.iloc[:750]
test_df = df_part2.iloc[750:1000] # Week 4: Recommended for evaluation
```

```
In [375...]: import pandas as pd
from collections import defaultdict
from sklearn.feature_extraction.text import TfidfVectorizer
import re

def tsv_to_keyword_dict(df):
    keywords = (
        df.iloc[:, 1]
        .str.lower()
        .str.replace(r"[.,!?:;]", "", regex=True)
        .str.split()
    )
    return dict(zip(df.iloc[:, 0].str.strip(), keywords))

user1_df = pd.read_csv('./data/user1.tsv', sep='\t')
user2_df = pd.read_csv('./data/user2.tsv', sep='\t')

user1_keywords = tsv_to_keyword_dict(user1_df)
user2_keywords = tsv_to_keyword_dict(user2_df)
```

```
In [376...]: user1_keywords
```

```
Out[376...]: {'dark': ['fire', 'enemy', 'pain', 'storm', 'fight'],
'sadness': ['cry', 'alone', 'heartbroken', 'tears', 'regret'],
'personal': ['dream', 'truth', 'life', 'growth', 'identity'],
'lifestyle': ['party', 'city', 'night', 'light', 'rhythm'],
'emotion': ['love', 'memory', 'hug', 'kiss', 'feel']}
```

```
In [377...]: user2_keywords
```

```
Out[377...]: {'sadness': ['lost', 'sorrow', 'goodbye', 'tears', 'silence'],
'emotion': ['romance', 'touch', 'feeling', 'kiss', 'memory']}
```

```
In [378...]: def get_liked_songs(train_df, user_keywords):
    liked = defaultdict(list)
    for _, row in train_df.iterrows():
        topic = row['predicted_topic']
        text = row['processed'] # already preprocessed (predict 1-3 week)
        if any(kw in text for kw in user_keywords.get(topic, [])):
            liked[topic].append(text)
    return liked
```

```
In [379...]: from sklearn.feature_extraction.text import TfidfVectorizer

def build_user_profile(liked_docs_by_topic):
    profiles = {}
    for topic, docs in liked_docs_by_topic.items():
        combined = " ".join(docs)
        vectorizer = TfidfVectorizer(max_features=500)
        tfidf = vectorizer.fit_transform([combined])
        tfidf_array = tfidf.toarray().flatten()
        feature_names = vectorizer.get_feature_names_out()
        top_indices = tfidf_array.argsort()[-20:][::-1]
        top_words = [(feature_names[i], tfidf_array[i]) for i in top_indices]
```

```
    profiles[topic] = top_words
return profiles
```

```
In [380...]:  
def print_user_profile(profile, user_name):  
    print(f"\n===== {user_name} Profile Keywords =====")  
    for topic, words in profile.items():  
        print(f"\nTopic: {topic}")  
        for word, score in words:  
            print(f"{word}: {score:.4f}")
```

```
In [381...]:  
# User 1  
user1_liked = get_liked_songs(train_df, user1_keywords)  
user1_profile = build_user_profile(user1_liked)  
print_user_profile(user1_profile, "User 1")  
  
# User 2  
user2_liked = get_liked_songs(train_df, user2_keywords)  
user2_profile = build_user_profile(user2_liked)  
print_user_profile(user2_profile, "User 2")  
  
# User 3 (custom)  
user3_keywords = {  
    'emotion': ['heart', 'love', 'soul'],  
    'lifestyle': ['party', 'car', 'drink'],  
    'personal': ['life', 'family', 'journey'],  
    'sadness': ['tears', 'cry', 'alone'],  
    'dark': ['blood', 'devil', 'death']  
}  
  
user3_liked = get_liked_songs(train_df, user3_keywords)  
user3_profile = build_user_profile(user3_liked)  
print_user_profile(user3_profile, "User 3")
```

===== User 1 Profile Keywords =====

Topic: dark  
fight: 0.3237  
know: 0.2857  
like: 0.2478  
come: 0.2165  
na: 0.1830  
black: 0.1719  
blood: 0.1652  
stand: 0.1607  
grind: 0.1496  
tell: 0.1473  
yeah: 0.1339  
time: 0.1317  
gon: 0.1317  
hand: 0.1272  
cause: 0.1138  
kill: 0.1138  
head: 0.1094  
leave: 0.1005  
life: 0.0960  
people: 0.0960

Topic: emotion  
good: 0.6710  
feel: 0.4217  
touch: 0.2678  
know: 0.2154  
hold: 0.2062  
want: 0.1262  
go: 0.1170  
kiss: 0.0923  
time: 0.0923  
like: 0.0923  
miss: 0.0893  
baby: 0.0893  
cause: 0.0862  
na: 0.0862  
vision: 0.0800  
heart: 0.0800  
video: 0.0739  
feelin: 0.0708  
loove: 0.0708  
lip: 0.0677

Topic: personal  
life: 0.5989  
live: 0.3215  
know: 0.2294  
na: 0.1815  
world: 0.1737  
change: 0.1543  
yeah: 0.1478  
time: 0.1452  
like: 0.1387  
dream: 0.1387  
come: 0.1361  
wan: 0.1141  
thing: 0.1050

think: 0.0959  
need: 0.0946  
thank: 0.0830  
teach: 0.0817  
lord: 0.0804  
want: 0.0804  
cause: 0.0804

Topic: lifestyle  
night: 0.3277  
come: 0.2798  
time: 0.2747  
song: 0.2647  
tonight: 0.2445  
home: 0.1966  
na: 0.1941  
long: 0.1916  
mind: 0.1916  
right: 0.1890  
sing: 0.1815  
like: 0.1790  
yeah: 0.1487  
wan: 0.1412  
know: 0.1361  
wait: 0.1260  
closer: 0.1185  
stranger: 0.1159  
play: 0.1134  
baby: 0.1109

Topic: sadness  
cry: 0.5163  
tear: 0.3526  
know: 0.3022  
steal: 0.2519  
club: 0.2519  
think: 0.1889  
mean: 0.1889  
baby: 0.1763  
leave: 0.1259  
smile: 0.1133  
heart: 0.1133  
true: 0.1133  
want: 0.1133  
word: 0.1007  
like: 0.1007  
say: 0.0881  
stay: 0.0881  
feel: 0.0756  
hold: 0.0756  
music: 0.0756

===== User 2 Profile Keywords =====

Topic: emotion  
good: 0.5463  
touch: 0.5281  
hold: 0.2368  
kiss: 0.1821  
know: 0.1639

vision: 0.1578  
video: 0.1457  
loove: 0.1396  
time: 0.1396  
feelin: 0.1336  
feel: 0.1275  
luck: 0.1214  
look: 0.1153  
lovin: 0.1153  
lip: 0.1093  
gim: 0.1032  
cause: 0.1032  
like: 0.0789  
go: 0.0789  
wait: 0.0728

Topic: sadness  
heart: 0.4397  
break: 0.3916  
away: 0.3023  
like: 0.1924  
fall: 0.1649  
know: 0.1649  
leave: 0.1580  
come: 0.1305  
na: 0.1305  
hard: 0.1305  
fade: 0.1305  
go: 0.1237  
feel: 0.1168  
think: 0.1168  
scar: 0.1168  
blame: 0.1099  
open: 0.1099  
time: 0.1031  
tear: 0.1031  
wan: 0.0962

===== User 3 Profile Keywords =====

Topic: dark  
blood: 0.3830  
know: 0.2214  
come: 0.2064  
like: 0.2064  
devil: 0.1990  
feel: 0.1915  
time: 0.1716  
death: 0.1691  
black: 0.1368  
hand: 0.1318  
stand: 0.1219  
cause: 0.1144  
fight: 0.1119  
pull: 0.1069  
evil: 0.1045  
cold: 0.1045  
hell: 0.0995  
na: 0.0995  
mind: 0.0995

head: 0.0970

Topic: personal  
life: 0.6555  
live: 0.3065  
know: 0.2086  
na: 0.1674  
change: 0.1518  
yeah: 0.1504  
world: 0.1461  
time: 0.1419  
like: 0.1362  
thing: 0.1078  
wan: 0.1007  
come: 0.0979  
think: 0.0951  
want: 0.0951  
need: 0.0894  
teach: 0.0865  
lord: 0.0823  
cause: 0.0823  
right: 0.0780  
go: 0.0780

Topic: lifestyle

song: 0.5485  
come: 0.3033  
home: 0.2618  
like: 0.2493  
night: 0.2077  
na: 0.1911  
wan: 0.1787  
sing: 0.1745  
country: 0.1579  
tonight: 0.1496  
yeah: 0.1496  
tire: 0.1454  
mind: 0.1330  
change: 0.1288  
spoil: 0.1246  
time: 0.1163  
play: 0.1122  
wait: 0.0873  
know: 0.0873  
cause: 0.0831

Topic: sadness

cry: 0.5543  
tear: 0.3785  
know: 0.3109  
club: 0.2704  
steal: 0.2704  
baby: 0.1893  
mean: 0.1893  
smile: 0.1217  
true: 0.1217  
like: 0.1081  
heart: 0.1081  
say: 0.0946  
stay: 0.0946

```
feel: 0.0811
music: 0.0811
forever: 0.0676
time: 0.0676
word: 0.0676
lift: 0.0541
wish: 0.0541
```

```
Topic: emotion
good: 0.6301
go: 0.3820
hold: 0.3572
feel: 0.2084
heart: 0.1737
know: 0.1588
miss: 0.1439
vision: 0.1290
video: 0.1191
want: 0.1191
love: 0.1092
vibe: 0.1042
hand: 0.1042
na: 0.1042
lovin: 0.0943
baby: 0.0893
gim: 0.0843
wait: 0.0794
time: 0.0794
right: 0.0794
```

To simulate a content-based recommender system, we utilized the best-performing topic classifier from Part 1 — Logistic Regression with CountVectorizer (max\_features=500) and lemmatization. The steps were as follows:

#### Topic Prediction:

For each song, we predicted its topic using the trained classifier.

#### User Interests Matching:

For each user, we assumed they "liked" any song from Weeks 1–3 (training data) that: Was predicted to be in a particular topic, and contained at least one keyword from the user's provided interest list (case-insensitive matching).

**For user Profile Construction:** All liked songs for each topic were concatenated into a single document. We applied TF-IDF vectorization (max\_features=500) to generate a representation of user interest in that topic. For each topic, we extracted the top 20 words with the highest TF-IDF scores to represent the user's profile.

#### Results & Interpretation:

**User 1 Dark:** High TF-IDF scores for fight, blood, kill, night indicate strong alignment with violent or intense themes.

**Emotion:** Terms like feel, touch, kiss, hold show emotional engagement.

**Sadness:** Words like regret, trust, blame, leave reflect melancholic tones.

**Lifestyle:** Includes musical context with words like sing, song, radio.

**Conclusion:** The user profile meaningfully captures emotional and thematic interests.

User 2. Sadness: Strong representation with words like heart, break, leave, crash.

Emotion: Dominated by soft emotional language such as lip, kiss, memory, feel.

Conclusion: Matches a gentle and melancholic user preference well.

User 3 (Custom) Emotion: Includes love, hold, vibe, feel, suggesting emotional resonance.

Sadness: Terms such as cry, tear, fear, steal express emotional depth.

Lifestyle: Includes song, country, tonight, night — typical of casual or musical interests.

Dark: Contains strong themes like blood, devil, death, hell.

Personal: life, world, change, time reflect general introspection.

Conclusion: The profile shows diverse and well-aligned interests across multiple themes.

## Q2.

We set M=20 to focus on the top 20 representative keywords per topic in the user profile, as these contribute most to the user's interest vector while reducing noise from low-weight terms.

In [382...]

```
def jaccard_similarity(user_vec, item_vecs):
    user_bin = (user_vec > 0).astype(int).flatten()
    item_bin = (item_vecs > 0).astype(int)

    sims = []
    for row in item_bin:
        row = row.toarray().flatten()
        intersection = np.logical_and(user_bin, row).sum()
        union = np.logical_or(user_bin, row).sum()
        sims.append(intersection / (union + 1e-10))
    return np.array(sims)

def dice_similarity(user_vec, item_vecs):
    user_bin = (user_vec > 0).astype(int).flatten()
    item_bin = (item_vecs > 0).astype(int)

    sims = []
    for row in item_bin:
        row_dense = row.toarray().flatten()
        intersection = np.sum(np.logical_and(user_bin, row_dense))
        denominator = np.sum(user_bin) + np.sum(row_dense)
        sims.append(2 * intersection / (denominator + 1e-10))
    return np.array(sims)
```

In [383...]

```
from sklearn.metrics.pairwise import cosine_similarity, euclidean_distances
import numpy as np

def compute_similarity(user_vec, item_vecs, metric='cosine'):
    if metric == 'cosine':
        return cosine_similarity(user_vec, item_vecs).flatten()
    elif metric == 'euclidean':
        return -euclidean_distances(user_vec, item_vecs).flatten()
    elif metric == 'jaccard':
        return jaccard_similarity(user_vec, item_vecs)
    elif metric == 'dice':
        return dice_similarity(user_vec, item_vecs)
    else:
        raise ValueError("Unsupported similarity metric")
```

In [384...]

```
import numpy as np
from sklearn.feature_extraction.text import TfidfVectorizer
from sklearn.metrics.pairwise import cosine_similarity

# Recommendation function: Recommend the top-N songs for a specific user in a ce
# M How many keywords before
def recommend_top_n(user_topic_profile, test_df, topic, N=10, M=20, metric='cosi
    user_keywords = dict(user_topic_profile[:M])
    vocab = list(user_keywords.keys())
    user_vector = np.array([user_keywords.get(word, 0.0) for word in vocab]).res
    topic_songs = test_df[test_df['predicted_topic'] == topic]

    if topic_songs.empty:
        return []

    vectorizer = TfidfVectorizer(vocabulary=vocab)
    test_tfidf = vectorizer.fit_transform(topic_songs['processed'])

    sims = compute_similarity(user_vector, test_tfidf, metric)

    top_n_indices = sims.argsort()[-N:][::-1]
    top_songs = topic_songs.iloc[top_n_indices].copy()
    top_songs['similarity'] = sims[top_n_indices]

    return top_songs[['track_name', 'artist_name', 'predicted_topic', 'similarit
```

In [385...]

```
def evaluate_user(user_keywords, user_profile, test_df, N=10, M=20, metric='cosi
    total_true_positives = 0
    total_recommended = 0
    total_relevant = 0

    for topic in user_keywords:
        if topic not in user_profile:
            continue
        recommended_df = recommend_top_n(user_profile[topic], test_df, topic, N=
        if recommended_df is None or recommended_df.empty:
            continue

        recommended_texts = recommended_df['track_name'].tolist()
        recommended_text_set = set(recommended_texts)

        relevant_keywords = set(user_keywords[topic])
        topic_test_songs = test_df[test_df['predicted_topic'] == topic]
        relevant_text_set = set()
        for _, row in topic_test_songs.iterrows():
            if any(kw in row['processed'] for kw in relevant_keywords):
                relevant_text_set.add(row['track_name'])

        true_positives = len(recommended_text_set & relevant_text_set)
        total_true_positives += true_positives
        total_recommended += len(recommended_text_set)
        total_relevant += len(relevant_text_set)

        precision = total_true_positives / total_recommended if total_recommended >
        recall = total_true_positives / total_relevant if total_relevant > 0 else 0.
        f1 = 2 * precision * recall / (precision + recall) if (precision + recall) e

    return precision, recall, f1
```

In [386...]

```
#Define four metrics which occur in Lecture
metrics = ['cosine', 'euclidean', 'jaccard', 'dice']

for metric in metrics:
    print(f"\n--- Using {metric.upper()} Similarity ---")
    for name, (keywords, profile) in {
        "User 1": (user1_keywords, user1_profile),
        "User 2": (user2_keywords, user2_profile),
        "User 3": (user3_keywords, user3_profile)
    }.items():
        p, r, f = evaluate_user(keywords, profile, test_df, N=10, M=20, metric=m)
        print(f"{name}: Precision@10 = {p:.4f}, Recall@10 = {r:.4f}, F1@10 = {f:.4f}")

--- Using COSINE Similarity ---
User 1: Precision@10 = 0.7400, Recall@10 = 0.4205, F1@10 = 0.5362
User 2: Precision@10 = 0.2500, Recall@10 = 0.5556, F1@10 = 0.3448
User 3: Precision@10 = 0.6000, Recall@10 = 0.4110, F1@10 = 0.4878

--- Using EUCLIDEAN Similarity ---
User 1: Precision@10 = 0.6800, Recall@10 = 0.3864, F1@10 = 0.4928
User 2: Precision@10 = 0.2500, Recall@10 = 0.5556, F1@10 = 0.3448
User 3: Precision@10 = 0.5000, Recall@10 = 0.3425, F1@10 = 0.4065

--- Using JACCARD Similarity ---
User 1: Precision@10 = 0.6200, Recall@10 = 0.3523, F1@10 = 0.4493
User 2: Precision@10 = 0.2500, Recall@10 = 0.5556, F1@10 = 0.3448
User 3: Precision@10 = 0.5000, Recall@10 = 0.3425, F1@10 = 0.4065

--- Using DICE Similarity ---
User 1: Precision@10 = 0.6200, Recall@10 = 0.3523, F1@10 = 0.4493
User 2: Precision@10 = 0.2500, Recall@10 = 0.5556, F1@10 = 0.3448
User 3: Precision@10 = 0.5000, Recall@10 = 0.3425, F1@10 = 0.4065
```

We compared four similarity metrics — Cosine, Euclidean, Jaccard, and Dice — for content-based song recommendation. Cosine similarity achieved the best balance of precision and recall, especially for emotionally rich profiles like User 1 and User 3.

In this recommendation system, we choose to use Precision@10 as the main evaluation indicator instead of Recall or F1-score for the following reasons:

1. The number of recommendations is limited, and precision is more important. In practical applications, the system usually only shows a limited number of recommendation results to users (for example, 10 songs). At this time, the "accuracy" of the recommendation is more critical than the "comprehensiveness", and the precision rate can better reflect whether the recommendation results are in line with the user's interests.
2. Ground truth is incomplete, and Recall is easily distorted. We construct "favorite songs" by matching user keywords, but this does not mean that we know "all" songs that the user likes. Therefore, the denominator of Recall may be underestimated or misestimated, resulting in a lack of authenticity in the calculation of Recall and F1.
3. F1-score and other methods weigh precision and recall, but we value "high-quality recommendations". F1-score will emphasize compensation in terms of recall, but this

system emphasizes a small number of accurate recommendations. Therefore, F1-score cannot accurately express the goal of the recommendation system.

The recommendation effect of User 1 was the most ideal. The Precision@10 value exceeded 0.74 with Consine Similarity, indicating that the recommendation system successfully captured his interest in the keywords.

User 2 It is indicated that some of the recommended songs are liked by users, but the overall hit rate is relatively low. Possible reasons include that user interest tags are relatively scattered, keywords are sparse, or there is less content in the training set that matches their interests, resulting in limited model performance.

User 3 has higher Precision@10, indicating that the recommendation system only covers a portion of the actual preferences.

## Part3

Actually,user4 is my best friend, who study EE in UNSW. Based on the songs he selected each week, I picked out the key words, constructed his user profile, and then made recommendations. The following is my code and evaluation. And the result is actually judged based on his own choice.

In [392...]

```
import random

# Week 1-3
week1_df = train_df.iloc[0:250]
week2_df = train_df.iloc[250:500]
week3_df = train_df.iloc[500:750]

# N=10 recommend to user
N = 10
week1_sample = week1_df.sample(N, random_state=42)
week2_sample = week2_df.sample(N, random_state=43)
week3_sample = week3_df.sample(N, random_state=44)

# concat
user_candidate_df = pd.concat([week1_sample, week2_sample, week3_sample])
```

In [393...]

```
user_candidate_df
```

Out[393...]

	artist_name	track_name	release_date	genre	lyrics	topic	full_text
142	skool 77	vivo hip hop (live)	2017	hip hop	head camden sure raise toast patron saint waif...	dark	skool 77 viv hip hop (liv hip ho head camc
6	rebelution	trap door	2018	reggae	long long road occur look shortcut wanna caus...	dark	rebelutic trap do reggae lor long roa oc
97	alec benjamin	outrunning karma	2018	pop	outrun karma charmer bug larva follow colorado...	dark	al benjam outrunnir karma pc outri karm
60	phish	we are come to outlive our brains	2018	blues	overlap future pass pass vapor light liquid bl...	dark	phish we a come outlive o brains blu
112	madeleine Peyroux	shout sister shout	2016	jazz	shout sister shout hallelujah shout sister sho...	dark	madeleir peyroi shout sist shout ja sho
181	janiva magness	what i could do	2018	blues	pay dues truth leave leave leave forget silenc...	sadness	jani magne what i cou do blues p dues
197	eric ethridge	if you met me first	2018	country	strangest creature fail prey like vulture hand...	dark	eric ethridg if you m me fir count stra
184	imagine dragons	natural	2018	rock	hold line give give tell house come consequenc...	dark	imagin dragoi natural ro hold lir give g
9	eli young band	never land	2017	country	word yeah wreck roll lips high good get bottle...	sadness	eli your band nev land count word ye wi

	artist_name	track_name	release_date	genre	lyrics	topic	full_text
104	larkin Poe	john the revelator	2017	blues	tell write revelator tell write revelator tell...	lifestyle	larkin po john th revelat blues t write
307	damian marley	slave mill	2017	reggae	remember slavery time work dollar slave grind ...	dark	damia marley slav mill regga rememb slave
484	gregory porter	holding on	2016	jazz	weight shoulder think easier heart grow colder...	emotion	grego port holding o jazz weig shoulder
460	haken	1985	2016	jazz	stand hand direction misalign play role cast s...	dark	haken 198 jazz star har directio misalign
265	kings of leon	walls	2016	rock	leave tryin right fight fight point waste spac...	dark	kings of le walls ro leave try right fig
435	louis armstrong	life is so peculiar	2016	jazz	life peculiar write burkejimmy heusen note mar...	personal	lou armstron life is : peculiar ja life
387	magic giant	window	2017	rock	home wanna window dream fell asleep wake ecsta...	sadness	magic gian window ro home wan windc drea
348	white denim	there's a brain in my head	2016	blues	atwo aone child guess gonna turn wild take loo...	dark	white deni there's brain in n head blu e
260	radiohead	daydreaming	2016	rock	dreamers learn learn point return return late ...	personal	radiohead daydreamir ro dreme learn lea
472	christie huff	black and white	2019	country	time game fight hurt tear stay straight core r...	dark	christie hu black ar whi country tin gar

	artist_name	track_name	release_date	genre	lyrics	topic	full_text
429	vince staples	big fish	2017	hip hop	want know say want want late night ballin coun...	lifestyle	vince stapl big fish h hop wa know say w
625	nick hakim	roller skates	2017	jazz	think take grant go change believe leave wait ...	personal	nick haki roller skat jazz thi take gran
561	midland	at least you cried	2017	country	word lips sound good right things like heart b...	sadness	midland least yo cried count word lips :
551	the afters	well done	2018	rock	like pain go worry world fade away like moment...	sadness	the afte well doi rock lil pain & worry w
704	ty segall	i sing them	2019	blues	life mystery look inside speak talk pray sing ...	lifestyle	ty segal sing the blues li mystery loc
507	leon bridges	there she goes	2016	blues	go move fast like train somebody babe go baby ...	emotion	leon bridg there sl goes blu go mo fas:
650	john gurney	drink i think	2019	country	feel come alive take black break skin shake fo...	dark	john gurne drink i thi country fe come a
544	the cure	in between days (45 version)	2019	rock	yesterday felt like yesterday want walk away c...	sadness	the cure between days (< version) ro ye:
543	avail hollywood	sinking in quick sand	2018	blues	bear legend fate loom prevail weight past weig...	personal	av. hollywo sinking quick sa blues be

	artist_name	track_name	release_date	genre	lyrics	topic	full_te
660	grupo maximo grado	el hombre del equipo	2016	pop	sell feel fool go lead kind realise nice nice ...	dark	grup maxim grado hombre d equipo po se
716	iration	fly with me	2017	reggae	relax know high star moon constellations leave...	sadness	iration t with n reggae rel know hig sta

In [394...]

```

import pandas as pd
import random

# in order to construct the user_profile ,these song user give like
def sample_tracks(df, n, seed):
    sample_df = df.sample(n=n, random_state=seed).reset_index(drop=False)

    id_list = sample_df['index'].tolist()          # Original Line number
    name_list = sample_df['track_name'].tolist()    # List of song titles
    topic_list=sample_df["topic"].tolist()

    return sample_df.drop(columns='index'), id_list, name_list

# user_candidate_df It's the "samplable" DataFrame in week1-3
sample_df, id_list, name_list = sample_tracks(user_candidate_df, n=17, seed=42)

for idx, (track_id, name) in enumerate(zip(id_list, name_list), 1):
    print(f"{idx:02d}. id={track_id:<4} track_name={name} ")

```

01. id=543 track\_name=sinking in quick sand
02. id=387 track\_name=window
03. id=704 track\_name=i sing them
04. id=260 track\_name=daydreaming
05. id=9 track\_name=never land
06. id=104 track\_name=john the revelator
07. id=660 track\_name=el hombre del equipo
08. id=507 track\_name=there she goes
09. id=460 track\_name=1985
10. id=142 track\_name=vivo hip hop (live)
11. id=112 track\_name=shout sister shout
12. id=348 track\_name=there's a brain in my head
13. id=181 track\_name=what i could do
14. id=265 track\_name=walls
15. id=484 track\_name=holding on
16. id=551 track\_name=well done
17. id=6 track\_name=trap door

In [395...]

```
liked_subset = user_candidate_df.loc[id_list].copy()
```

In [396...]

```
train_df.head()
```

Out[396...]

	artist_name	track_name	release_date	genre	lyrics	topic	full_text	processed
0	loving	the not real lake	2016	rock	awake know go see time clear world mirror worl...	dark	loving the not real lake rock awake know go se...	loving real lake rock awake know go see time c...
1	incubus	into the summer	2019	rock	shouldn summer pretty build spill ready overfl...	lifestyle	incubus into the summer rock shouldn summer pr...	incubus summer rock summer pretty build spill ...
2	reignwolf	hardcore	2016	blues	lose deep catch breath think say try break wal...	sadness	reignwolf hardcore blues lose deep catch breat...	reignwolf hardcore blue lose deep catch breath...
3	tedeschi trucks band	anyhow	2016	blues	run bitter taste take rest feel anchor soul pl...	sadness	tedeschi trucks band anyhow blues run bitter t...	tedeschi truck band anyhow blue run bitter tas...
4	lukas nelson and promise of the real	if i started over	2017	blues	think think different set apart sober mind sym...	dark	lukas nelson and promise of the real if i star...	lukas nelson promise real started blue think t...

In [397...]

```
liked_docs_by_topic = (
    liked_subset
        .groupby('predicted_topic')['processed']
        .apply(list)
        .to_dict()
)
```

In [398...]

```
liked_docs_by_topic.keys()
```

Out[398...]

```
dict_keys(['dark', 'emotion', 'lifestyle', 'personal', 'sadness'])
```

In [410...]

```
def recommend_global_top_k(user_profile: dict,
                           test_df: pd.DataFrame,
                           k_per_topic: int = 10,
```

```

        k_final: int = 10,
        metric: str = "cosine",
        m_keywords: int = 20):
    all_recs = []

    for topic, profile in user_profile.items():
        recs = recommend_top_n(profile, test_df, topic,
                               N=k_per_topic, M=m_keywords,
                               metric=metric)
        if not recs.empty:
            recs = recs.copy()
            recs['topic'] = topic
            all_recs.append(recs)

    if not all_recs:
        return pd.DataFrame()

    # Merge and sort globally by similarity
    merged = pd.concat(all_recs, ignore_index=True)
    global_top_df = (merged
                      .sort_values("similarity", ascending=False)
                      .head(k_final)
                      .reset_index(drop=True))

    return global_top_df

```

## Recommendation Result

In [411]: `print(recommended_df[['track_name', 'artist_name', 'topic', 'similarity']])`

	track_name	artist_name \
881	boy in the bubble	alec benjamin
943	speechless (full)	naomi scott
873	amsterdam	nothing but thieves
885	the war is over	cage the elephant
838	missed connection	the head and the heart
865	scared of the dark (feat. xxxtentacion)	lil wayne
891	strike the root	thievery corporation
917	bella ciao	klischée
942	whiskey fever	dorothy
784	around the corner	rick braun
944	once in a while	timeflies
770	horsefly	dirty heads
860	close enough	brett young
925	mississippi kisses	leon bridges
810	so close	notd
786	summertime is in our hands	michael franti & spearhead
993	find you	robert glasper experiment
969	got it good	justin moore
846	so cold in the summer	taylor mcferrin
764	call me	walshy fire
876	21 summer	brothers osborne
895	i wouldn't mind	he is we
795	it's only right	wallows
967	first time again	jason aldean
857	kiss me	magic!
864	rethymno	digitalluc
1003	hayati	tinariwen
920	jazz thing	gang starr
1006	puff your cares away	gramatik
931	100%	trina
913	pressure under fire	gov't mule
874	john cougar, john deere, john 3:16	keith urban
990	living it up	damian marley
851	hard case	tedeschi trucks band
862	shame	jeff beck
796	sunday morning	axian
861	happiness	needtobreathe
875	live without limit	jahmiel
884	rocky road	alborosie
811	big blue	vampire weekend
834	let me down slowly	alec benjamin
889	stay with me	ayokay
836	blackout	sublime with rome
774	the revival	the dear hunter
868	i can't stop dreaming	soja
769	one call away	charlie puth
826	the greatest	six60
789	will you be mine	anita baker
814	let you love me	rita ora
853	our hearts (feat. lucie silvas)	randy houser

topic	similarity
881	dark 0.518519
943	dark 0.461538
873	dark 0.461538
885	dark 0.400000
838	dark 0.400000
865	dark 0.400000
891	dark 0.333333

```
917      dark    0.333333
942      dark    0.333333
784      dark    0.333333
944      emotion 0.518519
770      emotion 0.461538
860      emotion 0.400000
925      emotion 0.333333
810      emotion 0.333333
786      emotion 0.333333
993      emotion 0.260870
969      emotion 0.260870
846      emotion 0.260870
764      emotion 0.260870
876      lifestyle 0.400000
895      lifestyle 0.400000
795      lifestyle 0.333333
967      lifestyle 0.333333
857      lifestyle 0.333333
864      lifestyle 0.333333
1003     lifestyle 0.333333
920      lifestyle 0.333333
1006     lifestyle 0.260870
931      lifestyle 0.260870
913      personal 0.400000
874      personal 0.333333
990      personal 0.260870
851      personal 0.260870
862      personal 0.181818
796      personal 0.181818
861      personal 0.181818
875      personal 0.181818
884      personal 0.181818
811      personal 0.181818
834      sadness 0.571429
889      sadness 0.518519
836      sadness 0.518519
774      sadness 0.518519
868      sadness 0.461538
769      sadness 0.461538
826      sadness 0.461538
789      sadness 0.461538
814      sadness 0.461538
853      sadness 0.400000
```

```
In [412...]: def evaluate_manual_feedback(recommended_df, liked_tracks):
    recommended = set(recommended_df['track_name'])
    relevant = set(liked_tracks)
    true_positive = recommended & relevant

    precision = len(true_positive) / len(recommended) if recommended else 0
    recall = len(true_positive) / len(relevant) if relevant else 0
    f1 = 2 * precision * recall / (precision + recall) if (precision + recall) else 0

    return precision, recall, f1
```

```
In [413...]: # First, do the global Top-10
global_rec = recommend_global_top_k(
    user_profile=user4_profile,
    test_df=test_df,
    k_per_topic=10,    # Grab 10 for each topic
```

```

        k_final=10,      # Only 10 songs were given in the end
        metric="cosine"
    )

true_likes = [s.lower() for s in
              ['boy in the bubble','close enough','truly madly deeply',
               "i wouldn't mind",'we find love','crossfire / so into you']]

p,r,f1 = evaluate_manual_feedback(global_rec, true_likes)
print(p,r,f1)

```

0.5 0.833333333333334 0.625

1. Segment the songs that users like in Week 1–3 and remove stop words, same as V4 processing;
2. Use TfidfVectorizer(max\_features=1000) to calculate the TF-IDF of each song;
3. Aggregate the TF-IDF vectors of the songs that users like by topic to get the user profile of each topic,same as part2
4. Take the top M=20 keywords with the highest TF-IDF value under each topic as the user interest representation of the topic.
5. Continuation of Part 2 solution: 1. Cosine similarity; 2. Take 10 candidates for each topic; 3. Sort by similarity globally and keep the top 10.
6. Evaluation indicators: Precision@10, Recall@10, F1@10, consistent with Part 2

We used 17 songs liked by the user from Week 1–3 (out of 30) to build a user profile. The recommendation system then produced 10 songs in Week 4, evaluated using standard @10 metrics.

Indicator	Part 2	Part 3 (real)
Precision@10	<b>0.74 / 0.25 / 0.60</b> (User 1 / 2 / 3)	<b>0.50</b>
Recall@10	0.42 / 0.56 / 0.41	0.83
F1@10	0.54 / 0.34 / 0.49	0.63

Analysis:

1. Precision decreases vs. Recall increases: A real user only receives 10 songs at a time, but still likes 5 songs → Recall 0.83;
2. Part 2 Assumption is too optimistic: Part 2 assumes that "all songs containing keywords are liked", resulting in a high precision, such as (User 1=0.74); In fact, users are more picky, and long-tail noise is filtered out.
3. Algorithm is stable: Cosine similarity still maintains the best F1 in real click scenarios; Jaccard/Dice is affected by binarization and has a low hit rate.

User Feedback (Talk-Aloud Summary):

- “I really liked ‘Boy in the Bubble’ – the lyrics are very thoughtful, and I’m into Alec Benjamin’s style.”
- “Some songs sounded too slow or mellow for my taste, like ‘Find You’. I’d probably skip that.”
- “I recognized a couple of artists like Keith Urban, which made me more interested in clicking those tracks.”
- “Some recommendations felt a bit too similar – like a lot of sad songs in a row.”
- “It would be nice to see more energetic or upbeat songs mixed in.”

For System:

- “It would be cool if I could see a little explanation, like ‘You liked X, so here’s Y’. Right now it feels like a black box—I don’t know what the system is basing this on.”
- “I wanted to skip some songs immediately. If the system could learn not just what I like, but also what I dislike or skip quickly, I think recommendations would improve faster.”

In [ ]:

The user's acceptance of recommended content depends more on Precision. The goal of the recommendation system is to get users to click/like the recommended items, rather than to cover as many items as possible. If Recall is high but Precision is low, users will be distracted by a large number of irrelevant recommendations and have a poor experience. For example, if a user only likes 2 out of 10 songs recommended, although Recall may be good, the user will feel that the recommendation quality is poor.

1. Cosine can identify the representativeness of high-weight keywords (such as love, miss) for user interests, while Euclidean/Jaccard/Dice either ignores direction (Jaccard/Dice) or is affected by length (Euclidean).
2. In text recommendation, we often use TF-IDF vectors to represent user portraits and song content, which are often high-dimensional sparse vectors.

So, cosine similarity has better semantic differentiation ability and stability in text recommendation, and is particularly suitable for user portrait and lyrics vectorization representation scenarios.